



MONASH University

程序代写代做 CS 编程辅导

Information Technology

FIT5202 (Version IV – Sort and Group By)



Week 4a – Parallel Sort

WeChat: cstutorcs

Assignment Project Exam Help

algorithm distributed systems database

Email: tutorcs@163.com

systems **computation knowledge management**

QQ: 749389476

design e-business model data mining intelligent

distributed systems database software engineering

<https://tutorcs.com>

computation knowledge management analysis

程序代写代做 CS编程辅导

TANIAR
LEUNG
RAHAYU
GOEL

Wiley Series on Parallel and Distributed Computi



High Performance
Database Processing

Grid Databases

High Performance Parallel Database
Processing and Grid Databases



WILEY

WeChat: cstutorcs
Assignment Project Exam Help
Last Week

Email: tutorcs@163.com

QQ: 749389476

DAVID TANIAR, CLEMENT H.C. LEUNG,
WENNY RAHAYU, and SUSHANT GOEL

<https://tutorcs.com>

- 5.1 Join Operations
- 5.2 Serial Join Algorithms
- 5.3 Parallel Join Algorithms
- 5.4 Cost Models
- 5.5 Parallel Join Optimization
- 5.6 Summary
- 5.7 Bibliographical Notes
- 5.8 Exercises

Revision

程序代写代做 CS编程辅导

• Exercise 1 (FLUX Quiz)

- Parallel Join algorithms for
Local Join
- A. TRUE
- B. FALSE



consists of two major phases: Data Partitioning, and

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Revision

程序代写代做 CS编程辅导

• Exercise 2 (FLUX Quiz)

- Parallel Join algorithms for the following queries are:
- A. ROJA and DOJA
- B. DER
- C. OJSO
- E. only A and B
- F. A, B and C are correct



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

程序代写代做 CS编程辅导

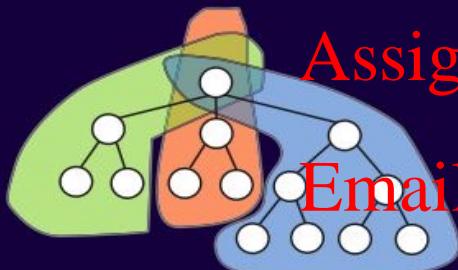
TANIA
LEUNG
RAHAYU
GOEL

Wiley Series on Parallel and Distributed Computi



High Performance
Database Processing
Grid Databases

WeChat: cstutorcs



Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

DAVID TANIA, CLEMENT H.C. LEUNG,
WENNY RAHAYU, and SUSHANT GOEL

<https://tutorcs.com>



WILEY

- 4.1 Sorting, Duplicate Removal and Aggregate
- 4.2 Serial External Sorting Method
- 4.3 Algorithms for Parallel External Sort
- 4.4 Parallel Algorithms for GroupBy Queries
- 4.5 Cost Models for Parallel Sort
- 4.6 Cost Models for Parallel GroupBy
- 4.7 Summary
- 4.8 Bibliographical Notes
- 4.9 Exercises

程序代写代做 CS 编程辅导

4.1. Sorting, and Serial Sorting



- Serial Sorting – INTERNAL
 - The data to be sorted fits entirely into the main memory

WeChat: cstutorcs

- Serial Sorting - EXTERNAL
 - The data to be sorted does NOT fit entirely into the main memory

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.1. Internal Serial Sorting (cont'd)



Bubble Sort

- Based on swapping
- It compares the first two elements, and if the first is greater than the second, it swaps them.
- It continues doing this for each pair of adjacent elements to the end of the data set.
- It then starts again with the first two elements, repeating until no swaps have occurred on the last pass.
- Example: 6 5 3 1 8 7 2 4

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

程序代写代做 CS 编程辅导

4.1. Internal Serial Sorting (cont'd)

Bubble Sort

6 5 3 1 8 7 2 4

5 6 3 1 8 7 2 4

5 3 6 1 8 7 2 4

5 3 1 6 8 7 2 4

5 3 1 6 8 7 2 4

5 3 1 6 7 8 2 4

5 3 1 6 7 2 8 4

5 3 1 6 7 2 4 8

5 3 1 6 7 2 4 8

3 5 1 6 7 2 4 8

3 1 5 6 7 2 4 8

3 1 5 6 7 2 4 8

3 1 5 6 7 2 4 8

3 1 5 6 7 2 4 8

3 1 5 6 2 7 4 8

3 1 5 6 2 4 8

QQ: 749389476

<https://tutorcs.com>



8 7 2 4

1 3 5 2 4 6 7 8

1 3 5 2 4 6 7 8

1 3 5 2 4 6 7 8

1 3 2 5 4 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

Finished

4.1. Internal Serial Sorting (cont'd)

- **Insertion Sort**

- Based on inserting
- It works by taking each element from the list one by one and inserting them in their correct position into a new sorted list. In arrays, the new list and the remaining elements can share the array's space, but insertion is expensive, requiring shifting all following elements over by one.
- Example: 6 5 3 1 8 7 2 4



3 1 8 7 2 4

WeChat: cstutorcs

Assignment Project Exam Help

Take out 6, and insert it in the previous list

Email: tutorcs@163.com

Take out 5, and insert it in the previous list

QQ: 749389476

Take out 3, and insert it in the previous list

Take out 1, and insert it in the previous list

https://tutorcs.com

Take out 8, and insert it in the previous list

Take out 7, and insert it in the previous list

Take out 2, and insert it in the previous list

Take out 4, and insert it in the previous list

6 5 3 1 8 7 2 4

5 6 3 1 8 7 2 4

3 5 6 1 8 7 2 4

1 3 5 6 8 7 2 4

1 3 5 6 8 7 2 4

1 3 5 6 7 8 2 4

1 2 3 5 6 7 8 4

1 2 3 4 5 6 7 8

Finished

4.1. Internal Serial Sorting (cont'd)



Quick Sort

- Quick Sort is a divide and conquer algorithm which relies on a partition operation: to partition an array an element called a *pivot* is selected.
- All elements smaller than the pivot are moved before it and all greater elements are moved after it.
- The lesser and greater sublists are then recursively sorted.
- The most complex issue in Quick Sort is choosing a good pivot element; consistently poor choices of pivots can result in drastically slower performance
- Example: 6 5 3 1 8 7 2 4

QQ: 749389476

<https://tutorcs.com> 6 5 3 1 8 7 2 4

4.1. Internal Serial Sorting (cont'd)



Quick Sort

- Quick Sort is a divide and conquer algorithm which relies on a partition operation: to partition an array an element called a *pivot* is selected.
- All elements smaller than the pivot are moved before it and all greater elements are moved after it.
- The lesser and greater sublists are then recursively sorted.
- The most complex issue in Quick Sort is choosing a good pivot element; consistently poor choices of pivots can result in drastically slower performance
- Example: 6 5 3 1 8 7 2 4

QQ: 749389476

<https://tutorcs.com> 6 5 3 1 8 7 2 4

Homework: Work out step-by-step
this Quick Sort example

程序代写代做 CS 编程辅导

4.2. Serial External Sorting



- Sorting is expressed by the ORDER BY clause in SQL
- Duplicate remove is identified by the keyword DISTINCT in SQL

Query 4.1:

```
Select *  
From STUDENT  
Order By Sdegree;
```

WeChat: cstutorcs

Query 4.3:

```
Select Distinct Sdegree  
From STUDENT;
```

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.2. Serial External Sorting (cont'd)

程序代写代做CS编程辅导



- External sorting assumes the data does not fit into main memory
- Most common external sorting is sort-merge
- Break the file up into unsorted subfiles, sort the subfiles, and then merge the subfiles into larger and larger sorted subfiles until the entire file is sorted

WeChat: cstutorcs
Algorithm: Serial External Sorting

// Sort phase - Pass 0
1. Read B pages at a time into memory
2. Sort them, and Write out a sub-file
3. Repeat steps 1-2 until all pages have been processed

// Merge phase - Pass 1, 2, ...
4. While the number of sub-files at end of previous pass is > 1
5. While there are sub-files to be merged from previous pass
6. Choose $B-1$ sorted sub-files from the previous pass
7. Read each sub-file into an input buffer page at a time
8. Merge these sub-files into one bigger sub-file
9. Write to the output buffer one page at a time

Email: tutorcs@163.com
QQ: 749389476
<https://tutorcs.com>

Figure 4.1 External sorting algorithm based on sort-merge

4.2. Serial External Sorting (cont'd)

程序代写代做CS编程辅导

- Example

- File size to be sorted = 108 pages, number of buffer (or memory size) = 5 pages
- Number of subfiles = 22 subfiles (the last subfile is only 3 pages long).
- Pass 0 (sorting phase): Sort each subfile, **read from disk, sort in main-memory**, and **write to disk** (Note: sorting the data in main-memory can use any fast in-memory sorting method, like Quick Sort)



WeChat: cstutorcs

- Merging phase: We use $B-1$ buffers (4 buffers) for input and 1 buffer for output
- Pass 1: Read 4 sorted subfiles and perform 4-way merging (apply a need k -way algorithm). Repeat the 4-way merging until all subfiles are processed. Result = 6 subfiles with 20 pages each (except the last one which has 8 pages)
- Pass 2: Repeat 4-way merging of the 6 subfiles like pass 1 above. Result = 2 subfiles
- Pass 3: Merge the last 2 subfiles
- Summary: 108 pages and 5 buffer pages require 4 passes

Assignment Project Exam Help

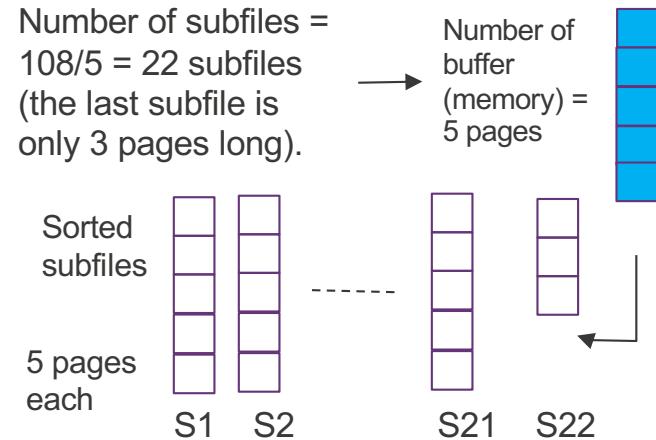
Email: tutorcs@163.com

QQ: 749389476

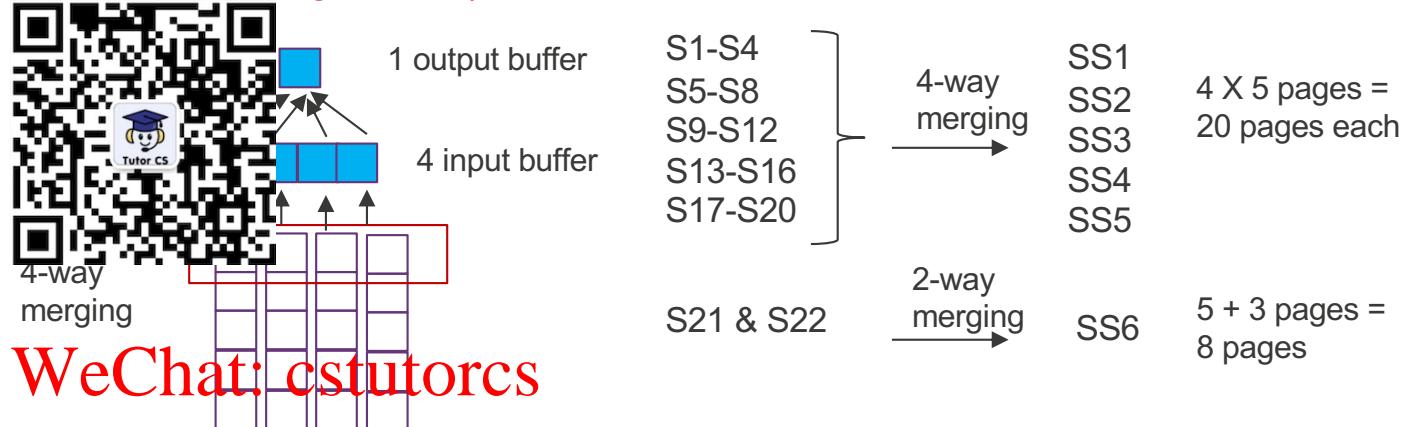
<https://tutorcs.com>

程序代写代做 CS编程辅导

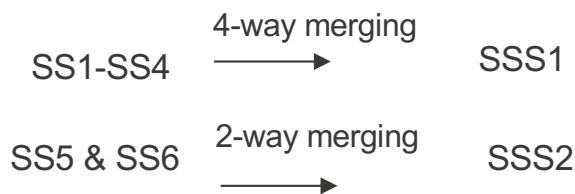
Pass 0 (sorting phase):



Pass 1 (Merging phase):



Pass 2 (Merging phase):



Pass 3 (Merging phase):

Email: tutorcs@163.com

4 X 20 pages = 80 pages

~~20 + 8 pages =~~

28 pages

QQ: 749389476

SSS1 & SSS2

2-way merging

SSSS1 (sorted array)

<https://tutorcs.com>

4.2. Serial External Sorting (cont'd)

程序代写代做CS编程辅导

Exercise 3 (FLUX Quiz)

- There are 150 data pages. The machine that we have has a limited memory, and can only take 8 pages at a time. How many passes will it take to sort the 150 data pages?



WeChat: cstutorcs

- A. 2
- B. 3
- C. 4
- D. 5

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

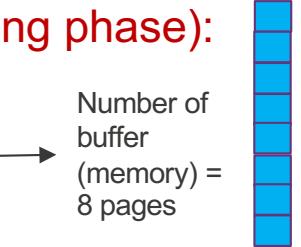
<https://tutorcs.com>

程序代写代做 CS编程辅导

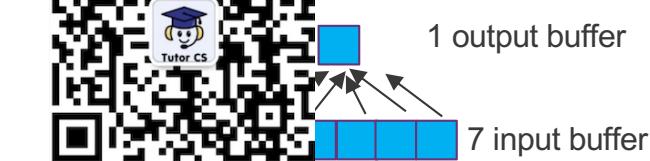
File size to be sorted = 150 pages, number of buffer (or memory size) = 8 pages

Pass 0 (sorting phase):

Number of subfiles =
 $150/8 = 19$ subfiles
(the last subfile is
only 6 pages long).



Pass 1 (Sorting phase):



S1-S7
S8-S14 } 7-way merging → SS1
SS2 7 X 8 pages = 56 pages each

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com
Pass 2 (Merging phase):

QQ: 749389476

SS1 & SS2 & SS3 → 3-way merging → SSS1 (sorted array)

<https://tutorcs.com>

4.2. Serial External Sorting (cont'd)

程序代写代做CS编程辅导

- Example



- Buffer size plays an important role in external sort

WeChat: cstutorcs

Table 4.1 Number of passes in serial external sorting as number of buffer increases

R	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1 million	20	10	7	5	3	3
10 million	23	12	8	6	4	3
100 million	26	14	9	7	4	4
1 billion	29	15	10	8	5	4



4.3. Parallel External Sort

5 different Algorithms

- Parallel Merge-All Sort
- Parallel Binary-Merge Sort
- Parallel Redistribution Binary-Merge Sort
- Parallel Redistribution Merge-All Sort
- Parallel Partitioned Sort



Without data
redistribution

With data
redistribution

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.3. Parallel External Sort (cont'd)



Parallel Merge-All

- A traditional approach
- Two phases: local sort and final merge
- Load balanced in local sort
- Problems with merging:
 - Heavy load on one processor
 - Network contention

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

WeChat: cstutorcs

Records from the child operator

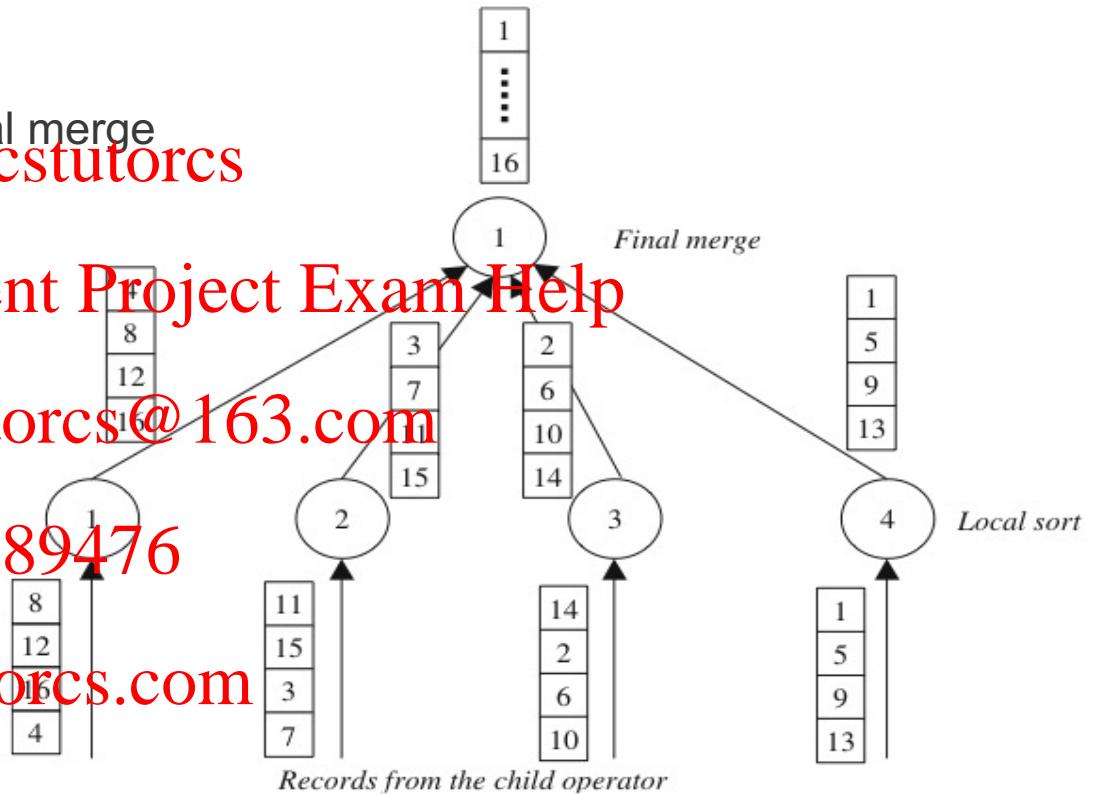


Figure 4.3 Parallel merge-all sort

4.3. Parallel External Sort (cont'd)



Parallel Binary-Merge Sort

- Local sort similar to traditional method
- Merging in pairs only
- Merging work is now spread to pipeline of processors, but merging is still heavy

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

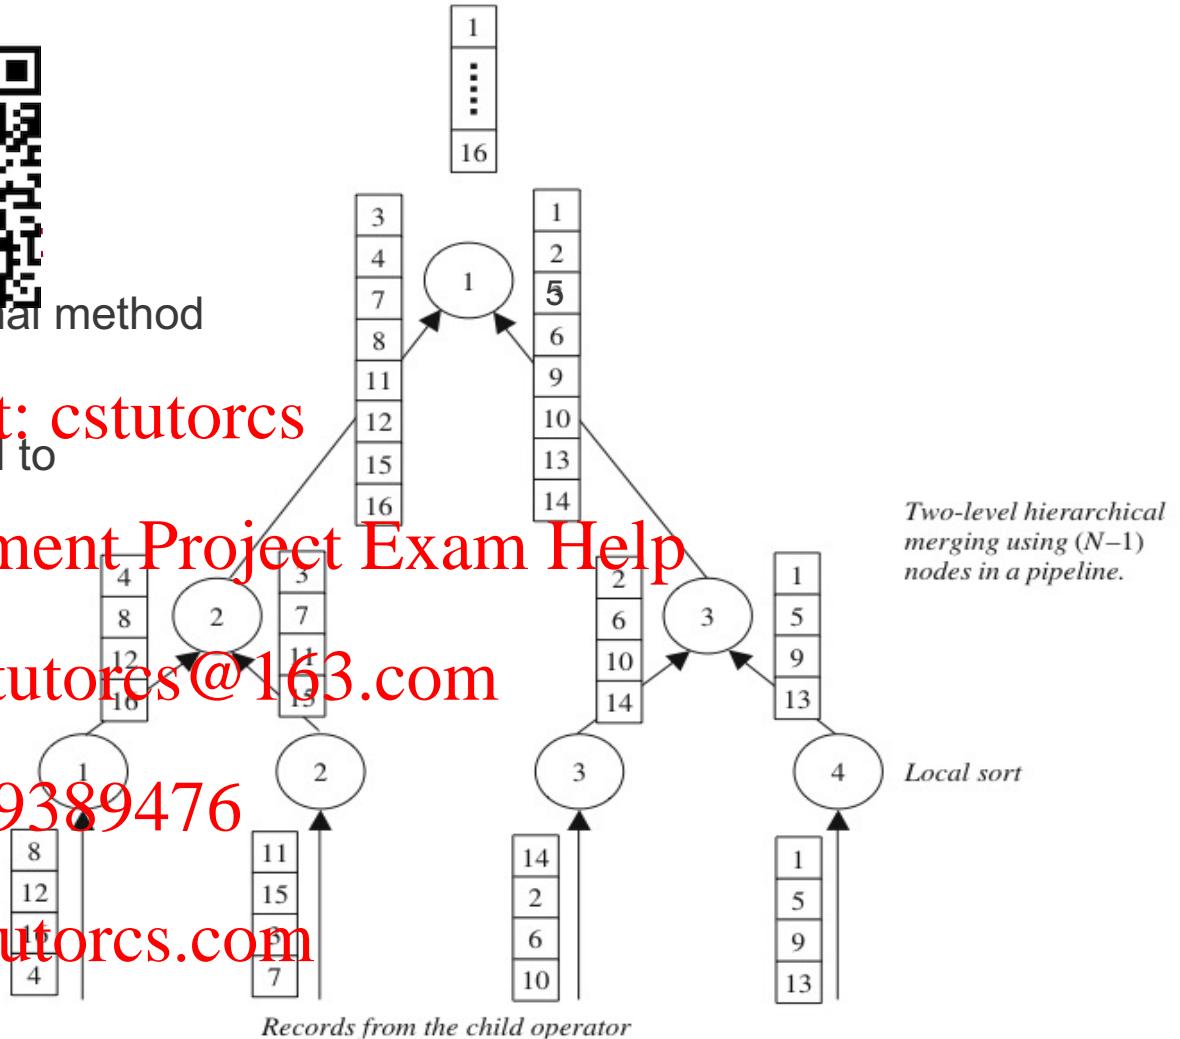


Figure 4.4 Parallel binary-merge sort

4.3. Parallel External Sort (cont'd)

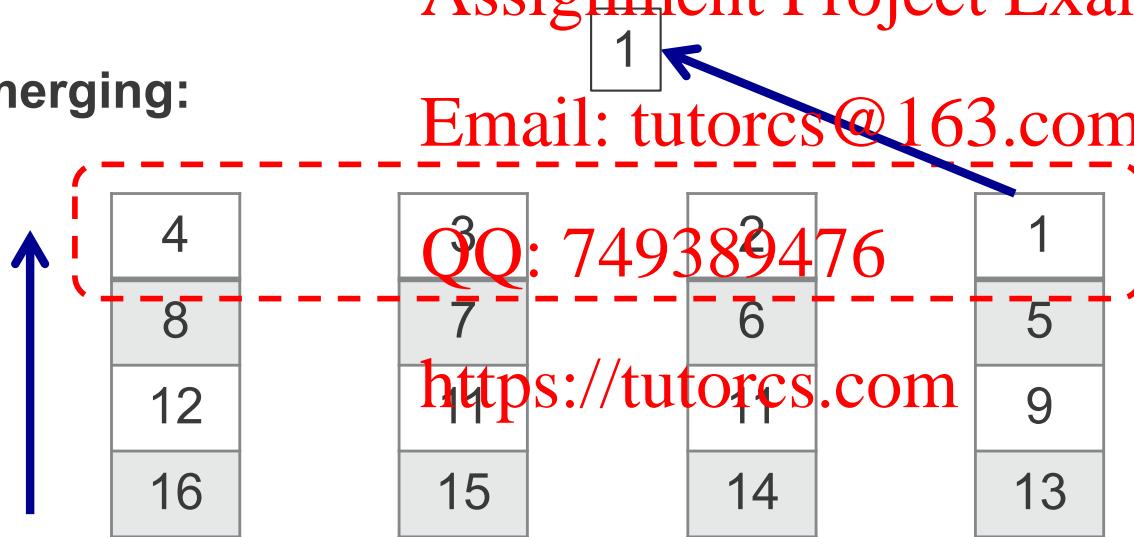
程序代写代做CS编程辅导

• Parallel Binary-Merge Sort

- Binary merging vs. k-way merging
- In k -way merging, searching for the smallest value among k partitions is done at the same time
- In binary merging, it is pairwise, but can be time consuming if the list is long
- System requirements: **k -way merging requires k files open simultaneously**, but the pipeline process in binary merging requires extra overheads

Assignment Project Exam Help

k -way merging:



Searching for the
smallest value
among the k values

4.3. Parallel External Sort (cont'd)



k-way merging:



Searching for the smallest value among the k values

4.3. Parallel External Sort (cont'd)

Parallel Binary-Merge Sort / Binary Merging step)

- Binary merging vs. k-way merging
- In **binary merging**, it is simple but can be time consuming if the list is long
- System requirements: the pipeline process in binary merging requires extra overheads



WeChat: cstutorcs

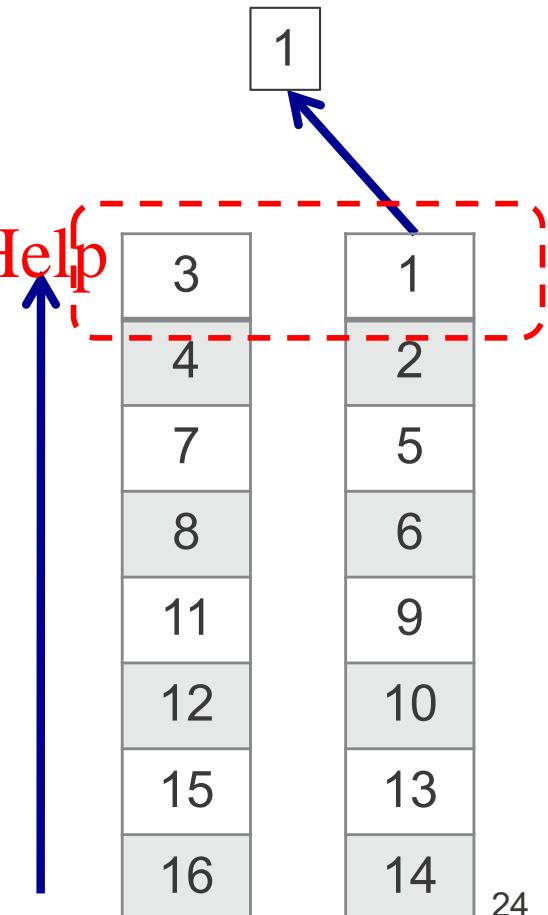
Assignment Project Exam Help

Binary merging:
Compare two values
only, but lists are
longer

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



4.3. Parallel External Sort (cont'd)

Parallel Binary-Merge Sort / Binary Merging step)

- Binary merging vs. k-way merging
- In **binary merging**, it is simple but can be time consuming if the list is long
- System requirements: the pipeline process in binary merging requires extra overheads



WeChat: cstutorcs

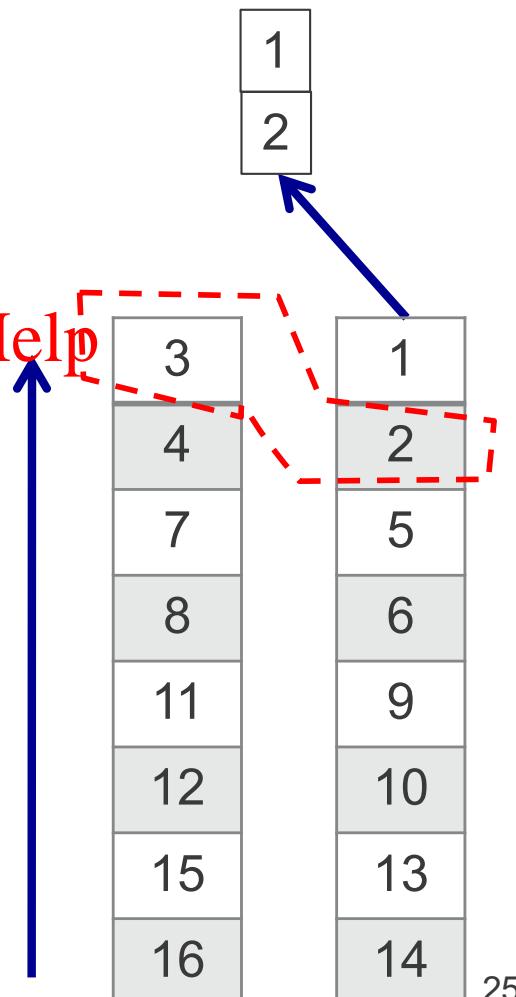
Assignment Project Exam Help

Binary merging:
Compare two values
only, but lists are
longer

Email: tutorcs@163.com

QQ: 749389476

And so on...
<https://tutorcs.com>



Parallel Redistribution Binary-Merge Sort

- Parallelism at all levels in the pipeline hierarchy
- Step 1: local sort
- Step 2: redistribute the results of local sort
- Step 3: merge using the same pool of processors

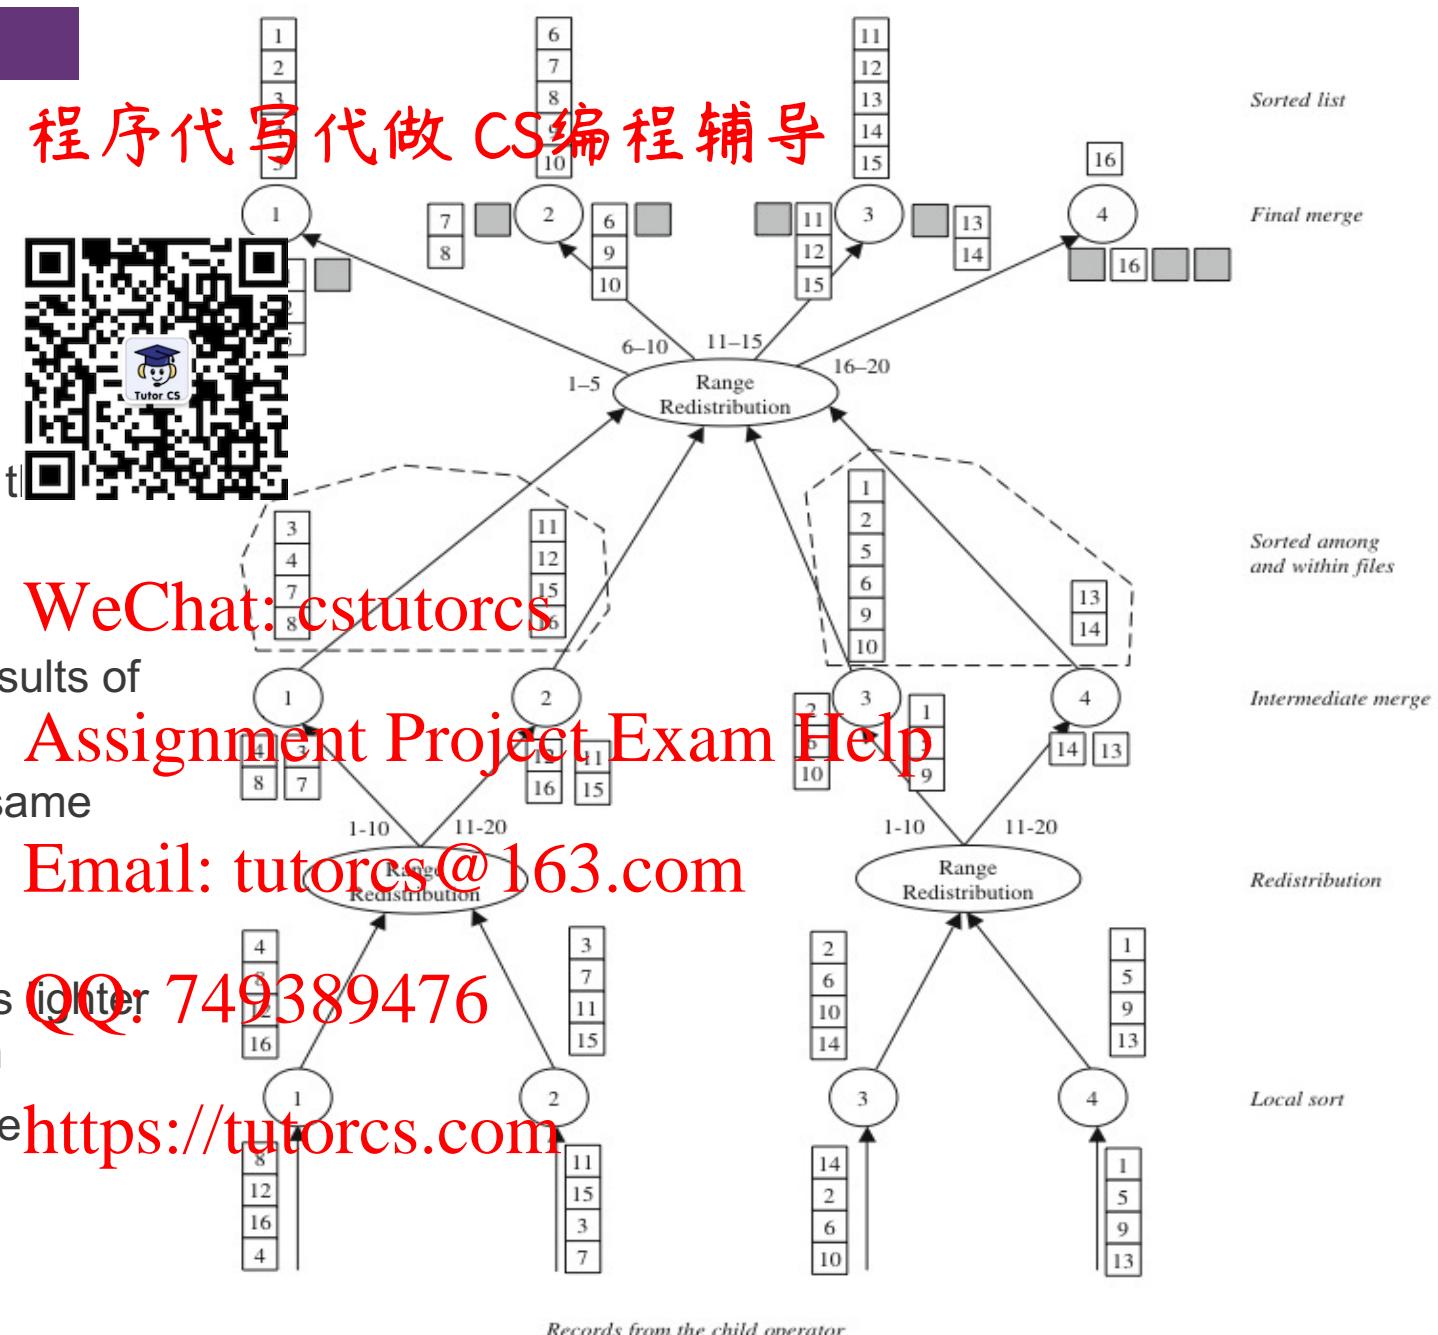


Figure 4.6 Parallel redistribution binary-merge sort

程序代写代做 CS编程辅导

- Parallel Redistribution Merge-All Sort

- Reduce the height of the tree and still maintain parallelism
- Like parallel merge-all sort, but with redistribution
- The advantage is true parallelism in merging
- Skew problem in the merging

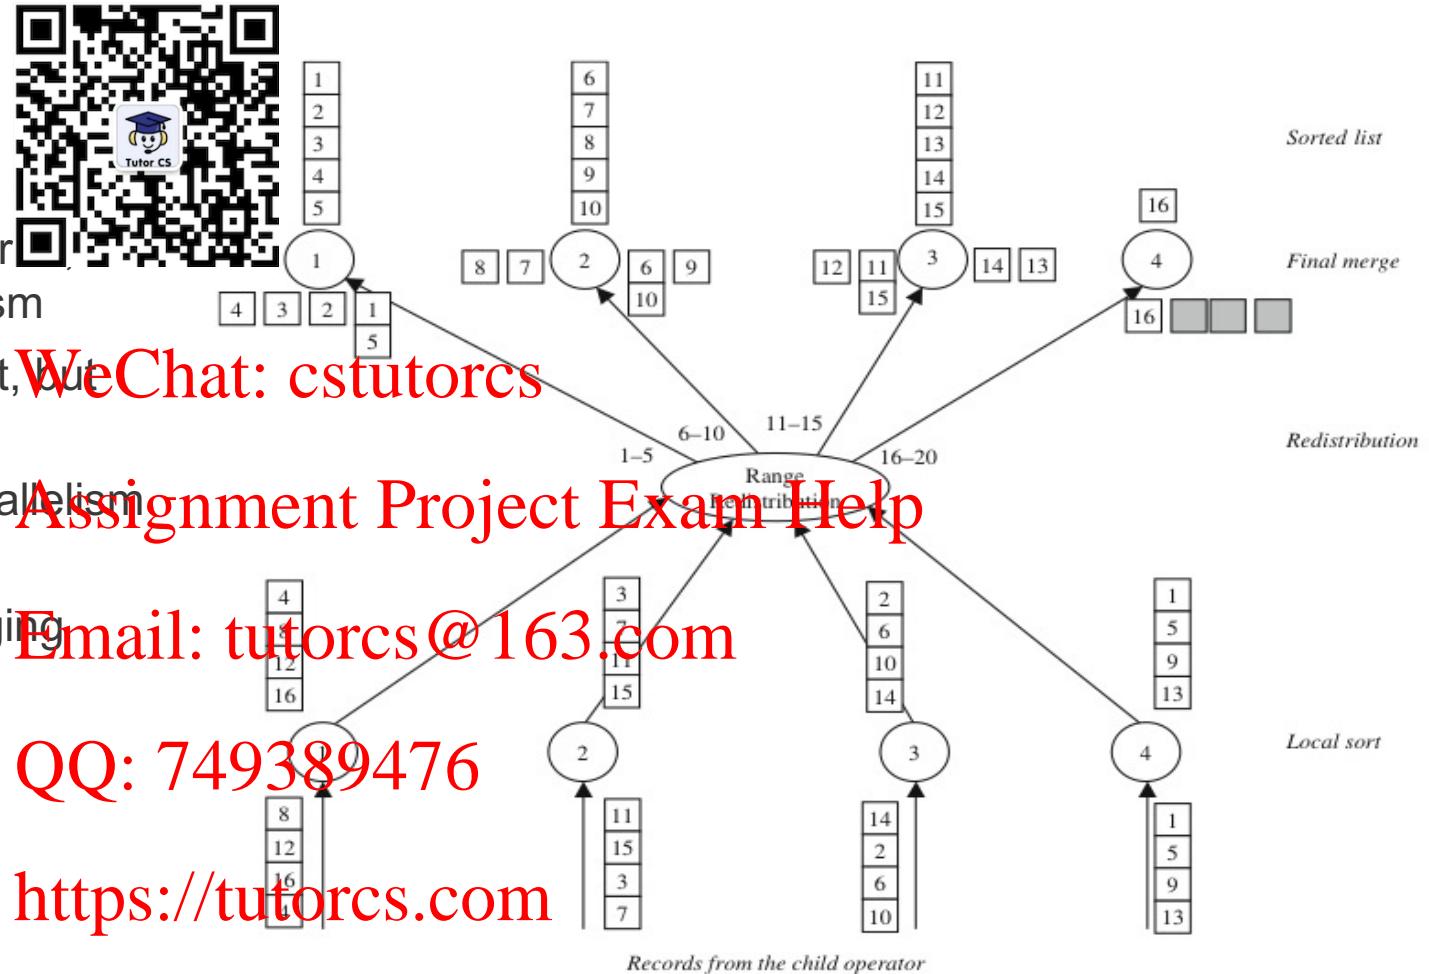


Figure 4.7 Parallel redistribution merge-all sort

程序代写代做 CS编程辅导

• Parallel Partitioned Sort

- Two stages: Partitioning step and Independent local work
- Partitioning (or range redistribution) may raise load skew
- Local sort is done after the partitioning, not before
- No merging is necessary
- Main problem: **Skew** produced by the partitioning

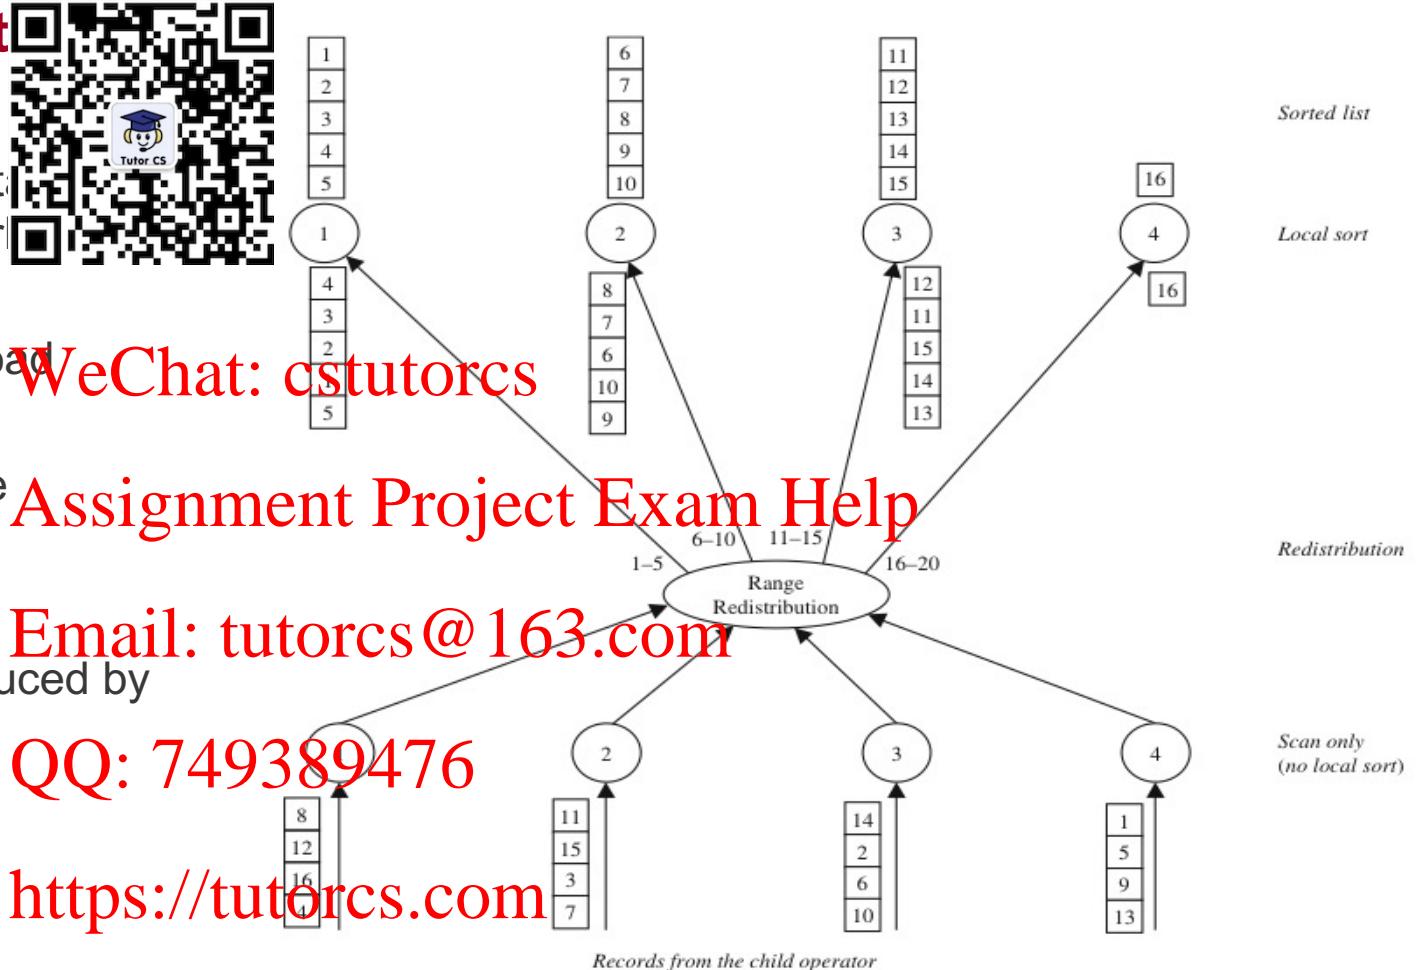


Figure 4.8 Parallel partitioned sort

4.2. Serial External Sorting (cont'd)

程序代写代做CS编程辅导

Exercise 4 (Home Work)

- Given a data set $D = \{55; 10; 4; 49; 90; 34; 76; 82; 56; 31; 25; 78; 56; 38; 32; 88; 9; 44; 98; 11; 70; 66; 89; 6\}$ and four processors, show step by step how the Parallel Partitioned Sort works.



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.2. Serial External Sorting (cont'd)

程序代写代做CS编程辅导

Exercise 5 (Difficult)

- Given the same dataset as the previous question, and 4 processors, show how **Load Balancing** is achieved in **Partitioned Sort**.



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



MONASH University

程序代写代做 CS 编程辅导

Information Technology

FIT5202 (Version IV – Sort and Group By)



Week 4b – Parallel Group By

WeChat: cstutorcs

Assignment Project Exam Help

algorithm distributed systems database

Email: tutorcs@163.com

systems **computation knowledge management**

QQ: 749389476

design e-business model data mining intelligent

distributed systems database software engineering

<https://tutorcs.com>

computation knowledge management analysis

程序代写代做 CS编程辅导

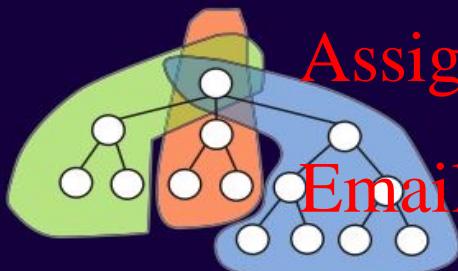
TANIA
R LEUNG
RAHAYU
GOEL

Wiley Series on Parallel and Distributed Computi



High Performance
Database Processing
Grid Databases

WeChat: cstutorcs



Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

DAVID TANIA, CLEMENT H.C. LEUNG,
WENNY RAHAYU, and SUSHANT GOEL

<https://tutorcs.com>



WILEY

- 4.1 Optimal Duplicate Removal and Aggregate
- 4.2 Serial External Sorting Method
- 4.3 Algorithms for Parallel External Sort
- 4.4 Parallel Algorithms for GroupBy Queries
- 4.5 Cost Models for Parallel Sort
- 4.6 Cost Models for Parallel GroupBy
- 4.7 Summary
- 4.8 Bibliographical Notes
- 4.9 Exercises

4.1. GroupBy, and Serial GroupBy

Select Suburb, Count(*)
From Student
Group By Suburb;

Student	Suburb
Adrian	Clayton
Beth	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(*)
From Student
Group By Suburb;

Processing Steps:

1. Read the first student record, and hash the suburb to the hash table

Student ID	Suburb
Adrian	Clayton
Bella	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

Hash the record using a certain hash function

Hash Table

1		
2		
3		
4		
5		
6		
7		
8	Clayton	1
9		

4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(*)
From Student
Group By Suburb;

Processing Steps:

1. Read the first student record, and hash the suburb to the hash table
2. Read the second record and hash it

Student ID	Suburb
Adrian	Clayton
Bella	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn



Hash Table

1	Hawthorn	1
2		
3		
4		
5		
6		
7		
8	Clayton	1
9		

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(*)
From Student
Group By Suburb;

Processing Steps:

1. Read the first student record, and hash the suburb to the hash table
2. Read the second record and hash it
3. Read the subsequent records one-by-one and hash them

Student ID	Suburb
Adrian	Clayton
Bella	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

Hash Table

1	Hawthorn	3
2	Caulfield	2
3	Malvern	1
4	Balwyn	1
5	Kew	1
6	Richmond	1
7	Elwood	1
8	Clayton	2
9	Doncaster	1

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(*)
From Student
Group By Suburb;

Processing Steps:

1. Read the first student record, and hash the suburb to the hash table
2. Read the second record and hash it
3. Read the subsequent records one-by-one and hash them
4. Read the Hash Table, and store this in disk as the query results



Hash Table in Main-Memory

1	Hawthorn	3
2	Caulfield	2
3	Malvern	1
4	Balwyn	1
5	Kew	1
6	Richmond	1
7	Elwood	1
8	Clayton	2
9	Doncaster	1

4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(*)
From Student
Group By Suburb;

Student	Suburb
Adrian	Clayton
Bethany	Hawthorn

This will work, if we assume that the main-memory can hold the entire Hash Table.

WeChat: cstutorcs

Assignment Project Exam Help
How about if the Hash Table is so big that it cannot fit into the main-memory.

Email: tutorcs@163.com

For example, how about if the main-memory can only hold 1 hash records at a time? How does the Group By processing work?

<https://tutorcs.com>

Leonard	Balwyn
Mary	Hawthorn

Hash Table

1	Hawthorn	3
2	Caulfield	2
3	Malvern	1
4	Balwyn	1
5	Kew	1
6	Richmond	1
7	Elwood	1
8	Clayton	2
9	Doncaster	1

4.1. Serial GroupBy Processing (cont'd)

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn



Hash Table

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Assume that the main-memory can hold 4 records in the hash table.

It needs a bigger hash table, but it doesn't have.

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

程序代写代做
CS编程辅导



based on the
Suburb
WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Ben	Hawthorn
DSpal	Caulfield
Garry	Hawthorn
Jessica	Caulfield
Mary	Hawthorn

Eric	Kew
Fred	Richmond
Katie	Malvern
Leonard	Balwyn

Adam	Clayton
Chris	Doncaster
Harold	Elwood
Irene	Clayton

In Main-Memory

In Disk

Ben	Hawthorn
Daniel	Caulfield
Garry	Hawthorn
Jessica	Caulfield
Mary	Hawthorn

Eric	Kew
Fred	Richmond
Katie	Malvern
Leonard	Balwyn

Adam	Clayton
Chris	Doncaster
Harold	Elwood
Irene	Clayton

Hash Processing

程序代写代做 CS编程辅导

In Main-Memory



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

Still in Disk

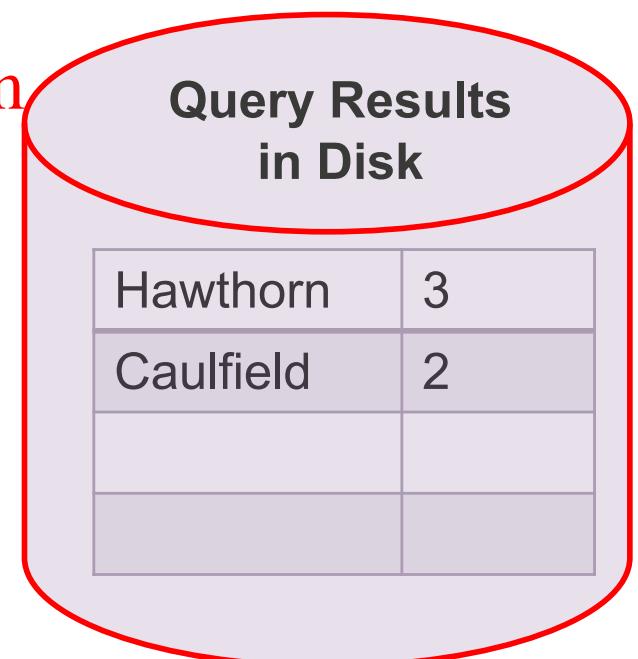
QQ: 749389476

<https://tutorcs.com>

Hash Table in Main-Memory

Hawthorn	3
Caulfield	2

Query Results
in Disk



Hash Processing

Hash Table in Main-Memory

Ben	Hawthorn
Daniel	Caulfield
Garry	Hawthorn
Jessica	Caulfield
Mary	Hawthorn

Eric	Kew
Fred	Richmond
Katie	Malvern
Leonard	Balwyn

Adam	Clayton
Chris	Doncaster
Harold	Elwood
Irene	Clayton

程序代写代做 CS 编程辅导

Flush to Disk



WeChat: cstutorcs

Assignment Project Exam Help
Load to Main-Memory

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>
Still in Disk

Kew	1
Malvern	1
Richmond	1
Balwyn	1

Query Results in Disk

Hawthorn	3
Caulfield	2
Kew	1
Caulfield	1
Richmond	1
Balwyn	1

Hash Processing

Hash Table in Main-Memory

Ben	Hawthorn
Daniel	Caulfield
Garry	Hawthorn
Jessica	Caulfield
Mary	Hawthorn

Eric	Kew
Fred	Richmond
Katie	Malvern
Leonard	Balwyn

Adam	Clayton
Chris	Doncaster
Harold	Elwood
Irene	Clayton

程序代写代做 CS 编程辅导

Flush to Disk



Flush to Disk

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

Load to Main-Memory

Clayton	2
Doncaster	1
Elwood	1

Hawthorn	3
Caulfield	2
Kew	1
Caulfield	1
Richmond	1
Balwyn	1
Clayton	2
Doncaster	1
Elwood	1

Query Results
in Disk

4.4. Parallel GroupBy



- Traditional methods (Shuttle Sort and Hierarchical Merging)
- Two-phase method
- Redistribution method

Without data redistribution

With data redistribution

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.4. Parallel GroupBy (cont'd)

程序代写代做CS编程辅导



• Traditional Method

- Step 1: local aggregate in each processor
- Step 2: global aggregation
- May use a Merge-All or Hierarchical method
- Need to pay a special attention to some aggregate functions (AVG) when performing a local aggregate process

WeChat: cstutorcs

Assignment Project Exam Help

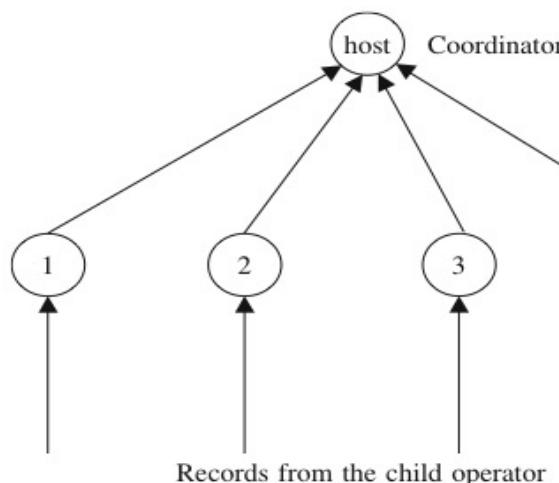


Figure 4.10 Traditional method

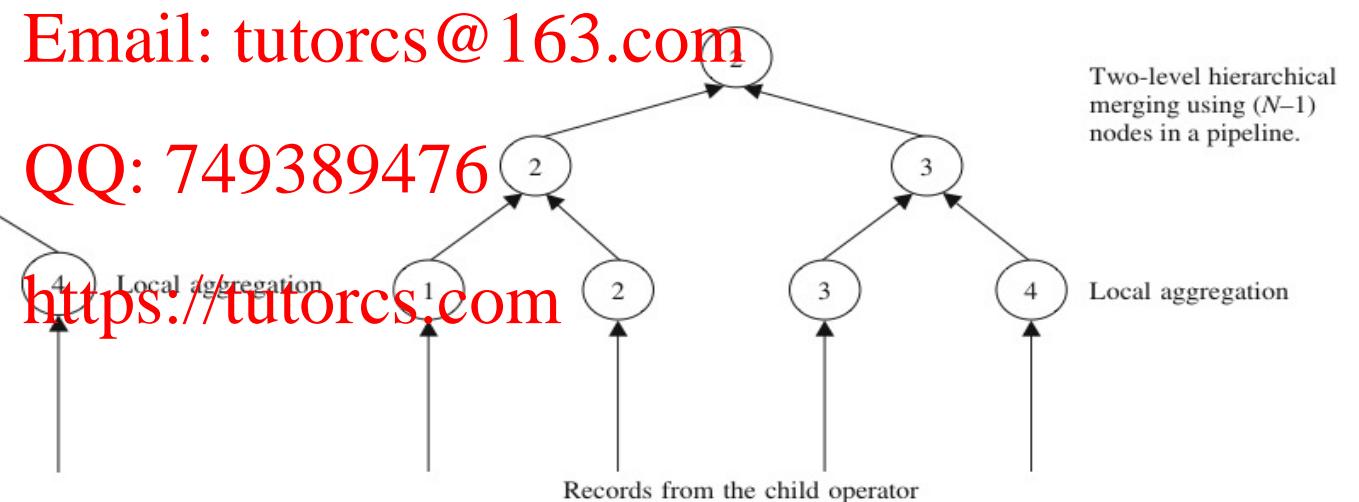
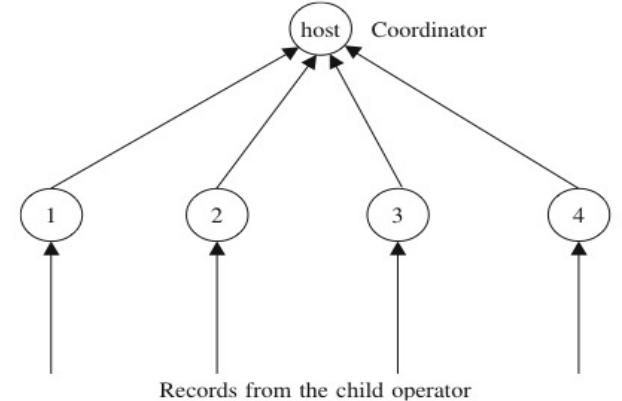


Figure 4.11 Hierarchical merging method

4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All

程序代写代做 CS编程辅导



WeChat: cstutorcs
Initial Data Placement
Assignment Project Exam Help

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Processor 2

Email: tutorcs@163.com

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All

程序代写代做 CS编程辅导

Clayton	2
Caulfield	1
Malvern	1
Vermont	1

Hawthorn	1
Richmond	1
Elwood	1
Malvern	1
Kew	1

Balwyn	1
Hawthorn	1
Kew	1
Richmond	1
Vermont	1

Elwood	1
Kew	1
Balwyn	1
Malvern	1
Clayton	1



Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont



Processor 2

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew



Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont



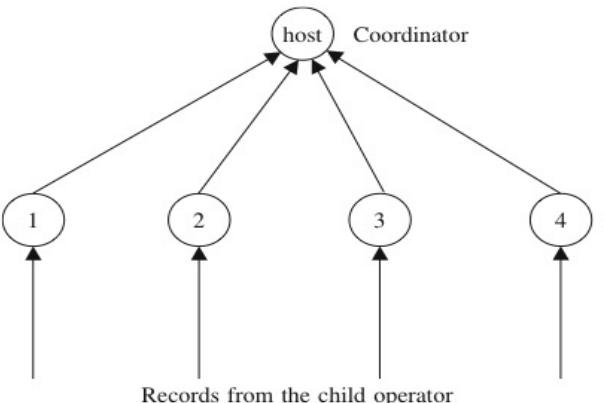
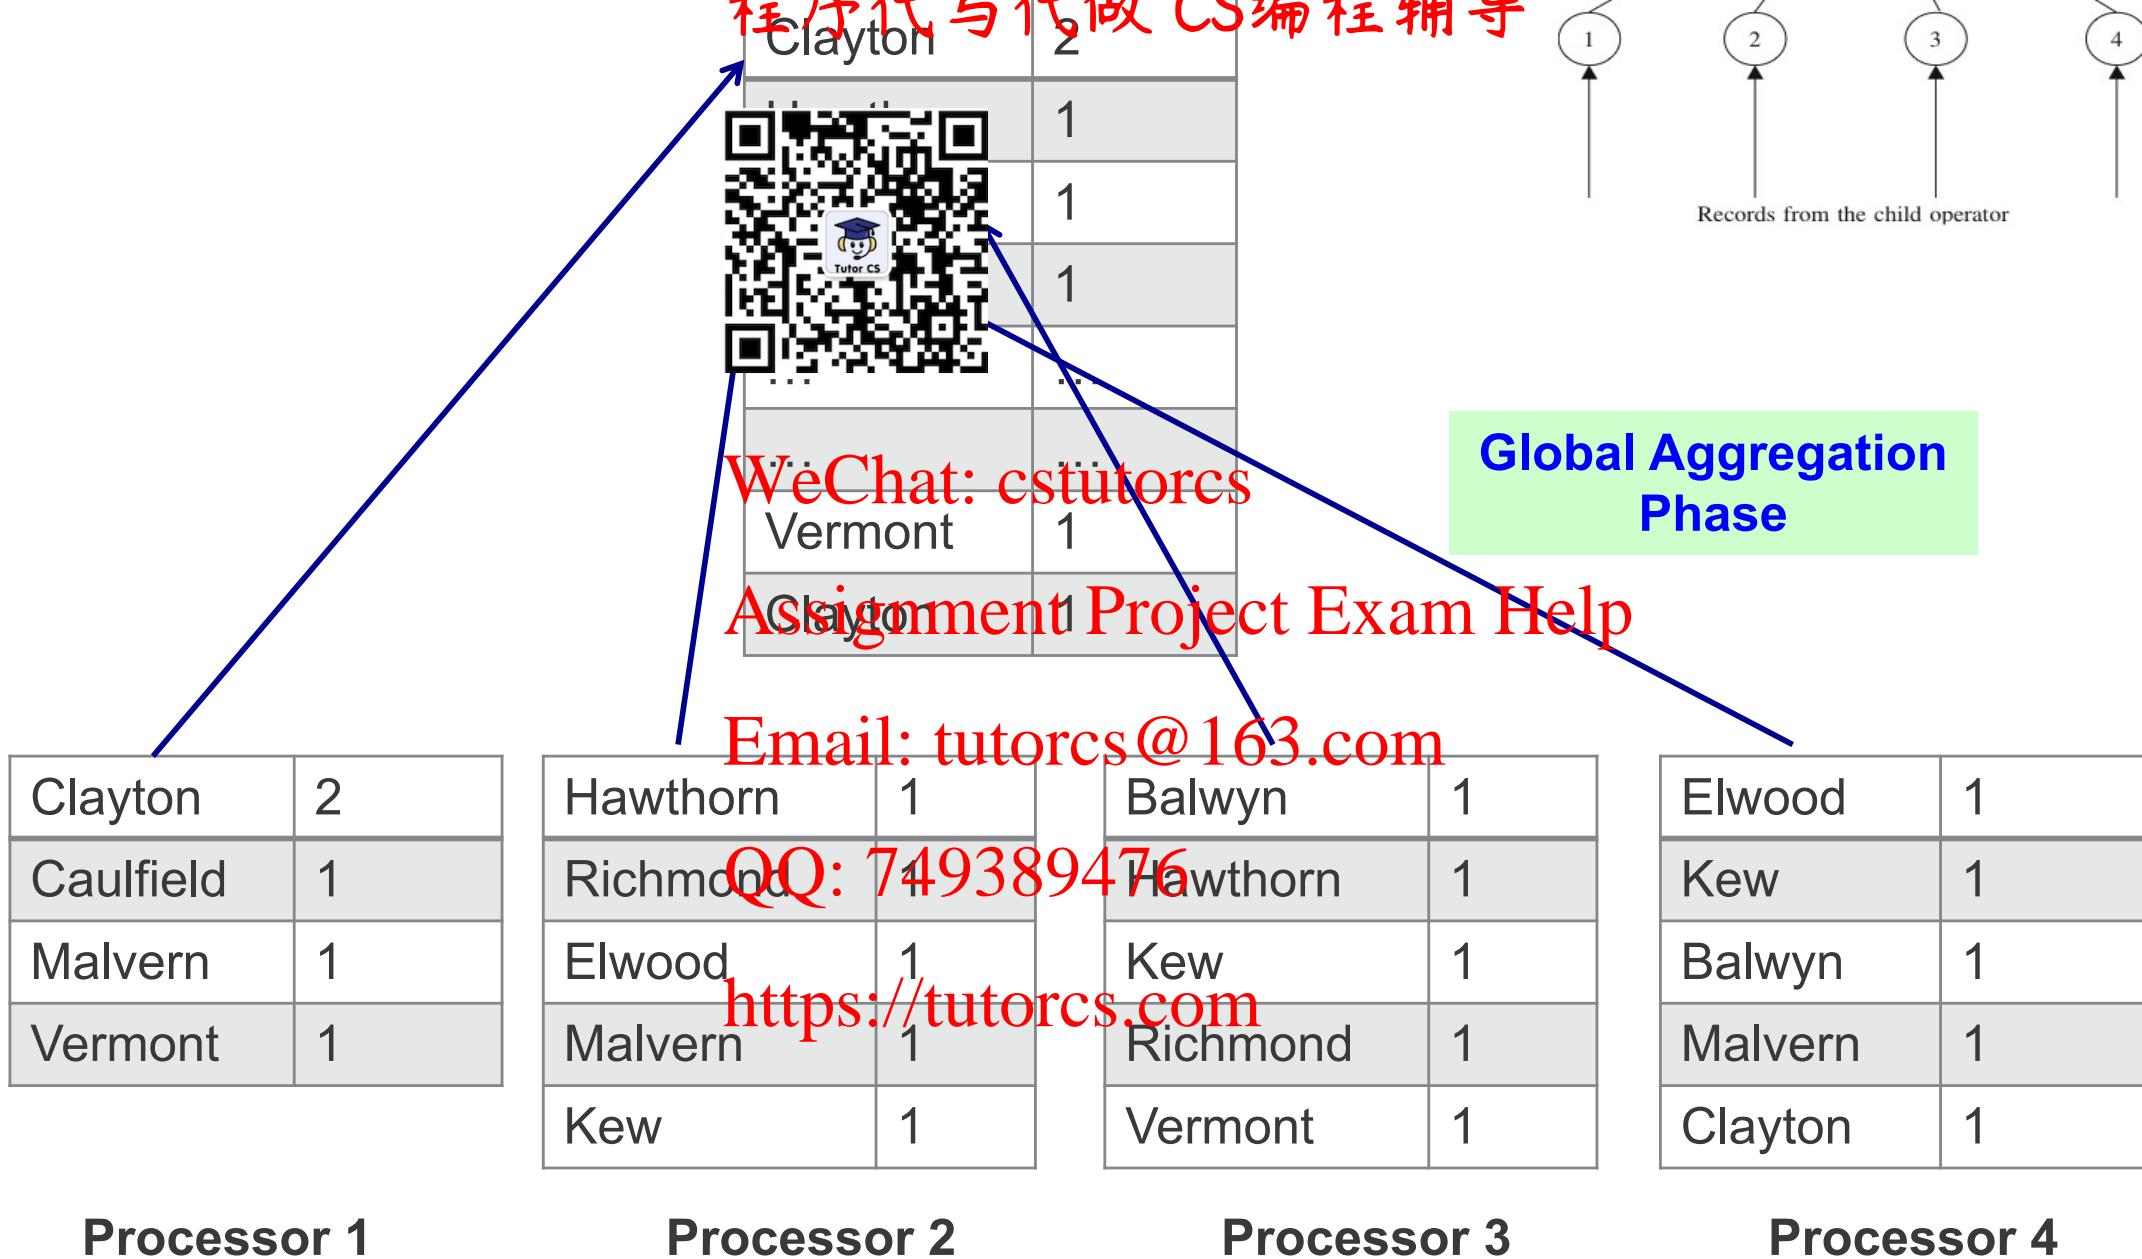
Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All

程序代写代做 CS 编程辅导



4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All

程序代写代做 CS 编程辅导

Clayton	3
	2
	2
	...
	...
Vermont	2

WeChat: cstutorcs

Assignment Project Exam Help

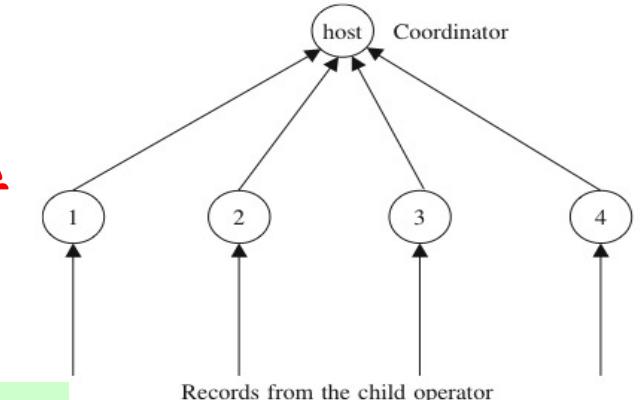
Clayton	2
Hawthorn	1
Balwyn	1
Elwood	1
...	...
...	...
Vermont	1
Clayton	1

QQ: 749389476

<https://tutorcs.com>

Email: tutorcs@163.com

Global Aggregation Phase



程序代写代做 CS编程辅导

• Exercise 6 (FLUX Quiz)

- The limitations of the Trapatch (Merge All) to process a Group By query are:

- A. Global aggregation is carried out by one processor
- B. Network bottleneck when sending the local aggregation results to the coordinator
- C. No parallelism in the global aggregation phase
- D. All of the above
- E. Some of the above

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.4. Parallel GroupBy (cont'd)

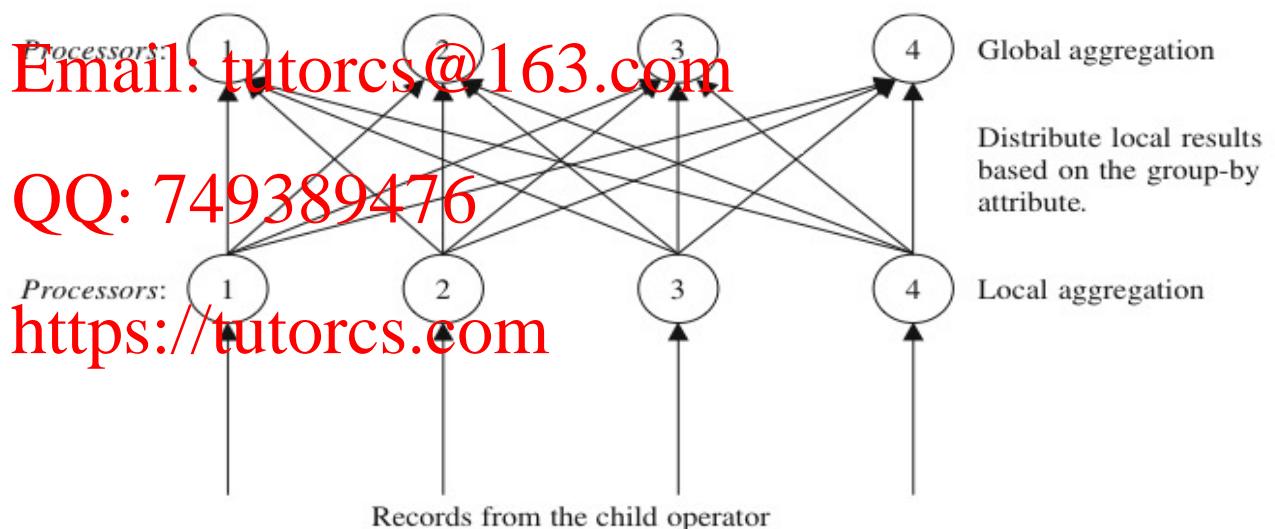
程序代写代做CS编程辅导



Two-Phase Method

- Step 1: local aggregate in each processor. Each processor groups local records according to the groupby attribute
- Step 2: global aggregation where all temp results from each processor are redistributed and then final aggregate is performed in each processor

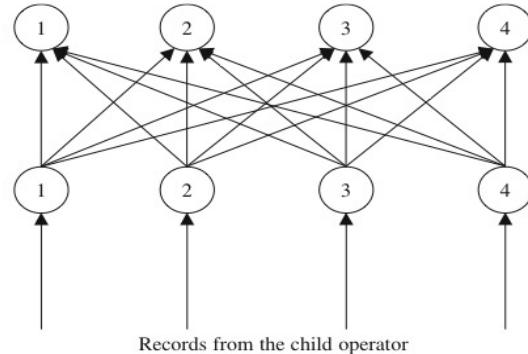
WeChat: cstutorcs
Assignment Project Exam Help



4.4. Parallel GroupBy (cont'd)

- Two-Phase Method

程序代写代做 CS编程辅导



WeChat: cstutorcs
Initial Data Placement
Assignment Project Exam Help

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Processor 2

Email: tutorcs@163.com
QQ: 749389476

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

4.4. Parallel GroupBy (cont'd)

- Two-Phase Method

程序代写代做 CS 编程辅导

Clayton	2
Caulfield	1
Malvern	1
Vermont	1

Hawthorn	1
Richmond	1
Elwood	1
Malvern	1
Kew	1

Balwyn	1
Hawthorn	1
Kew	1
Richmond	1
Vermont	1

Elwood	1
Kew	1
Balwyn	1
Malvern	1
Clayton	1



Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont



Processor 2 Processor 3

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew



Processor 4

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

4.4. Parallel GroupBy (cont'd)

- Two-Phase Method

Clayton	2
Balwyn	1
Elwood	1
Caulfield	1
Elwood	1
Balwyn	1
Clayton	1

程序代写代做 CS 编程辅导

Hawthorn	1
Hawthorn	1
Kew	1
Kew	1
Kew	1

Malvern	1
Malvern	1
Malvern	1

Richmond	1
Vermont	1
Richmond	1
Vermont	1

Clayton	2
Caulfield	1
Malvern	1
Vermont	1

WeChat: cstutorcs

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Hawthorn	1
Richmond	1
Elwood	1
Malvern	1
Kew	1

Balwyn	1
Hawthorn	1
Kew	1
Richmond	1
Vermont	1

Elwood	1
Kew	1
Balwyn	1
Malvern	1
Clayton	1

Processor 1

Processor 2

Processor 3

Processor 4

4.4. Parallel GroupBy (cont'd)

- Two-Phase Method

程序代写代做CS编程辅导

Final Results

Clayton	3
Balwyn	2
Elwood	2
Caulfield	1

Hawthorn	1
Kew	1

Malvern	3
---------	---

Richmond	2
Vermont	2

WeChat: cstutorcs

Clayton	2
Balwyn	1
Elwood	1
Caulfield	1
Elwood	1
Balwyn	1
Clayton	1

Hawthorn	1
Hawthorn	1
Kew	1
Kew	1
Kew	1

Malvern	1
Malvern	1
Malvern	1

Richmond	1
Vermont	1
Richmond	1
Vermont	1

Processor 1

Processor 2

Processor 3

Processor 4

QQ: 749389476

<https://tutorcs.com>

Global Aggregation Phase
Email: tutorcs@163.com

Final Results

4.4. Parallel GroupBy (cont'd)

程序代写代做CS编程辅导



• Redistribution Method

- Step 1 (Partitioning phase): redistribute raw records to all processors
- Step 2 (Aggregation phase): each processor performs a local aggregation

WeChat: cstutorcs

Assignment Project Exam Help

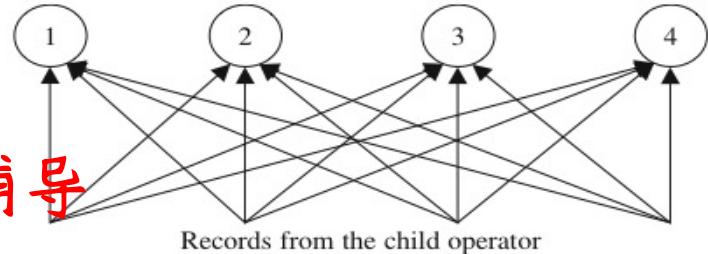


Figure 4.13 Redistribution method

4.4. Parallel GroupBy (cont'd)

- Redistribution Method

程序代写代做 CS编程辅导



WeChat: cstutorcs
Initial Data Placement
Assignment Project Exam Help

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Processor 2

Email: tutorcs@163.com
QQ: 749389476

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

4.4. Parallel GroupBy (cont'd)

- Redistribution Method

Adam	Clayton
Kelly	Balwyn
Peter	Elwood
Ben	Clayton
Chris	Caulfield
Harold	Elwood
Roger	Balwyn
Tracy	Clayton

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

程序代写代做 CS编程辅导

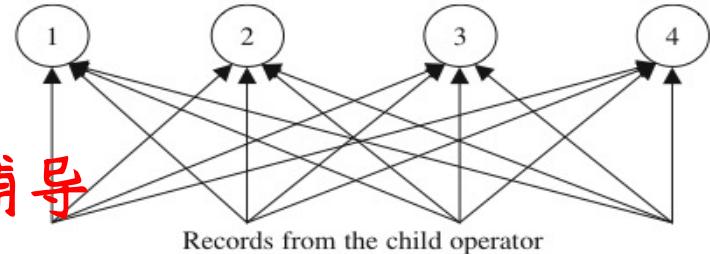
Partitioning Phase



Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

WeChat: cstutorcs
Assignment Project Exam Help

Email: tutorcs@163.com



George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

QQ: 749389476

<https://tutorcs.com>

4.4. Parallel GroupBy (cont'd)

· Redistribution Method

程序代写代做CS编程辅导

Final results

Clayton	3
Balwyn	2
Elwood	2
Caulfield	1

Hawthorn	2
Kew	2

Malvern	3
---------	---

Richmond	2
Vermont	2

Adam	Clayton
Kelly	Balwyn
Peter	Elwood
Ben	Clayton
Chris	Caulfield
Harold	Elwood
Roger	Balwyn
Tracy	Clayton

WeChat: cstutorcs

Assignment Project Exam Help

Aggregation Phase

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 1

Processor 2

Processor 3

Processor 4

4.4. Parallel GroupBy (cont'd)

• Redistribution Method

程序代写代做 CS编程辅导

Clayton	3
Balwyn	2
Elwood	2
Caulfield	1



Malvern	3
---------	---

Richmond	2
Vermont	2

Adam	Clayton
Kelly	Balwyn
Peter	Elwood
Ben	Clayton
Chris	Caulfield
Harold	Elwood
Roger	Balwyn
Tracy	Clayton

WeChat: cstutorcs
Assignment Project Exam Help
What is the problem here?

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 1

Processor 2

Processor 3

Processor 4

4.4. Parallel GroupBy (cont'd)

Redistribution Method (Task Stealing)

程序代写代做 CS 编程辅导

Create 5 buckets, instead of 4



Fred	
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Processor 2 tutorcs@163.com

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

QQ: 749389476

<https://tutorcs.com>

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

4.4. Parallel GroupBy (cont'd)

· Redistribution Method (Task Stealing)

程序代写代做 CS编程辅导



Adam	Clayton
Kelly	Balwyn
Ben	Clayton
Chris	Caulfield
Roger	Balwyn
Tracy	Clayton

Peter	Elwood
Harold	Elwood

Processor 1

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew

Processor 2

QQ: 749389476

<https://tutorcs.com>

Task stealing

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

Processor 3

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 4

4.4. Parallel GroupBy (cont'd)

· Redistribution Method (Task Stealing)

程序代写代做 CS 编程辅导

Clayton	3
Balwyn	2
Caulfield	1



Hawthorn	
Kew	



Malvern	3
Elwood	2

Richmond	2
Vermont	2



Adam	Clayton
Kelly	Balwyn
Ben	Clayton
Chris	Caulfield
Roger	Balwyn
Tracy	Clayton

Processor 1

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Processor 2

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Peter	Elwood
Harold	Elwood

Processor 3

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 4

程序代写代做 CS编程辅导

· Exercise 7 (FLUX Quiz)

- The Redistribution Method: balancing option, through the Task Stealing method.
- The Two-Phase Method: a load balancing problem.
- A. TRUE
- B. FALSE



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

4.7. Summary

程序代写代做 CS 编程辅导

- Parallel groupby algorithm



- Traditional methods (merging)
- Hierarchical methods
- Two-phase method - – Local aggregation **before** data redistribution
- Redistribution method - – Local aggregation **after** data redistribution

WeChat: cstutorcs

- Two-phase** and **Redistribution** methods perform better than the traditional and hierarchical merging methods

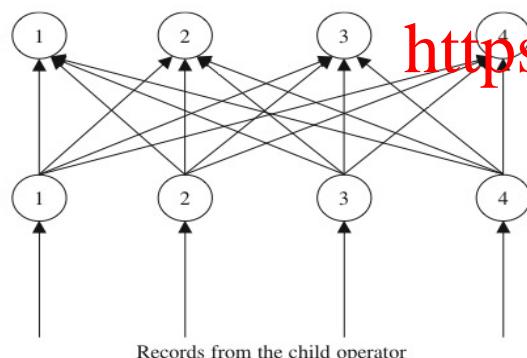
Assignment Project Exam Help

- Two-phase method** works well when the number of groups is small, whereas the **Redistribution method** works well when the number of groups is large

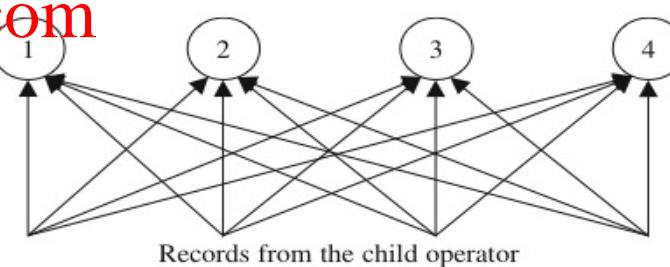
QQ: 749389476

Ambuj and Naughton. "Adaptive parallel aggregation algorithms." (1995):

Why??



<https://tutorcs.com>



程序代写代做 CS编程辅导

Homework Exercises

- 1. Show how Load Balancing Task Stealing be achieved in the Two Phase Method (using the same sample code) – EASY
- 2. Why is the Two-Phase Method good when the number of groups is small, whereas the Redistribution Method good when the number of groups is large? – MORE CHALLENGING

Assignment Project Exam Help

- 3. In what scenario may super linear speed up be achieved? – MORE CHALLENGING
- (Hints: See slides #8-#13 → the hash table cannot fit into main-memory)

QQ: 749389476

<https://tutorcs.com>

