



程序代写代做 CS 编程辅导



MONASH
INFORMATION
TECHNOLOGY

Introduction to Machine Learning

WeChat: cstutorcs

Assignment Project Exam Help

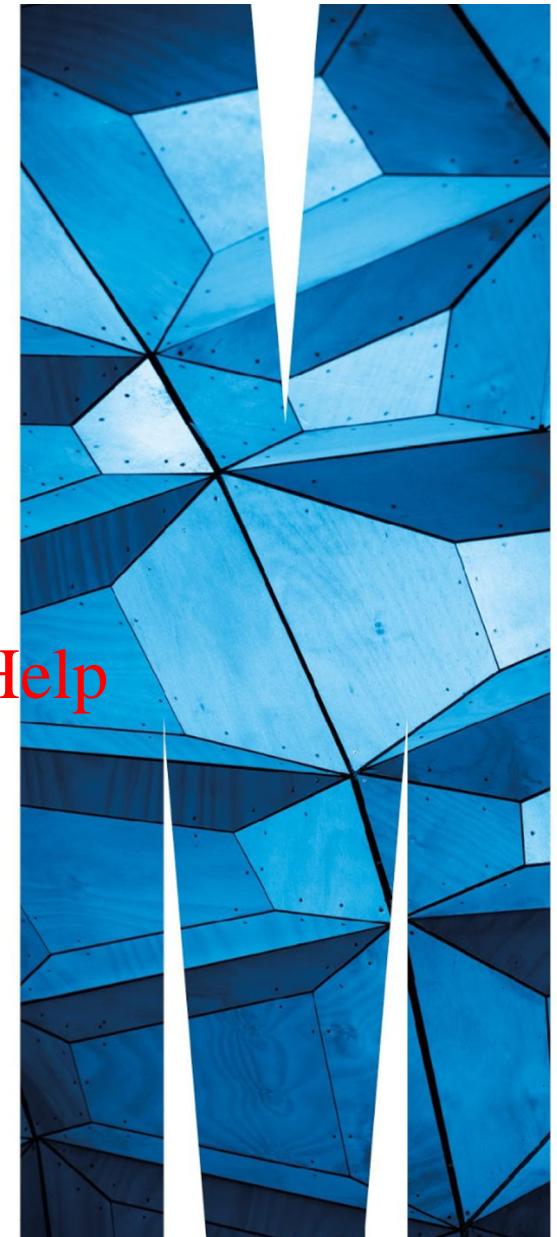
Developed by Prajwol Sangat

Updated by Chee-Ming Ting (3 April 2021)

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Last week

程序代写代做 CS编程辅导

Parallel Aggregation

Parallel Sort

Parallel Group-By



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

This week

程序代写代做 CS编程辅导

What is Machine Learning
Machine Learning Benefits
Types of Machine Learning
Feature Engineering



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

程序代写代做 CS编程辅导



According to [McKinsey study](#), 35% of what consumers purchase on Amazon and 75% of what they watch on Netflix is driven by machine learning-based product recommendations.

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

What is Machine Learning?

程序代写代做 CS编程辅导



What is Machine Learning?



A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E", (Tom Mitchell, 1997)

Face recognition



WeChat: cstutores

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Examples

程序代写代做 CS编程辅导



Detecting Spam Emails



Detect credit card fraud



MONASH
University



Task E
of
pairs

Task T
given an question,
find the best
answer

Performance P
how accurate the
answer is

WeChat: cstutorcs

Examples of spam
emails and not-
spam emails

To assign a label
“spam” or “not-
spam” to an email

how accurate spam
email can be
detected

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

Data collected for
credit-card
transactions deemed
as fraud and not-fraud

To assign a label
“fraud” or “not fraud”
to a given credit-card
transaction

how accurate a credit-
card fraud transaction
can be detected.

<https://tutorcs.com>

Elements of machine learning

① Data

feature $x \in \mathbb{R}^d$

label $y \in Y$

Dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^n$

Data processing

feature extraction,
feature selection,
feature transformation,
feature reduction,
feature scaling, feature
normalization



Predictive Model $Y = f_{\theta}(X)$

Assignment Project Exam Help
Model Learning (Training)

- Find an optimal model f_{θ} (by estimating model parameters θ) using training data
- Based on loss function (e.g., minimize error between true and predicted labels)

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

③ Assessment

How well is f_{θ} doing
w.r.t data \mathcal{D} ?

Model Testing $\hat{y} = f_{\theta}(x_{test})$

Test the learned model in predicting
unseen test data

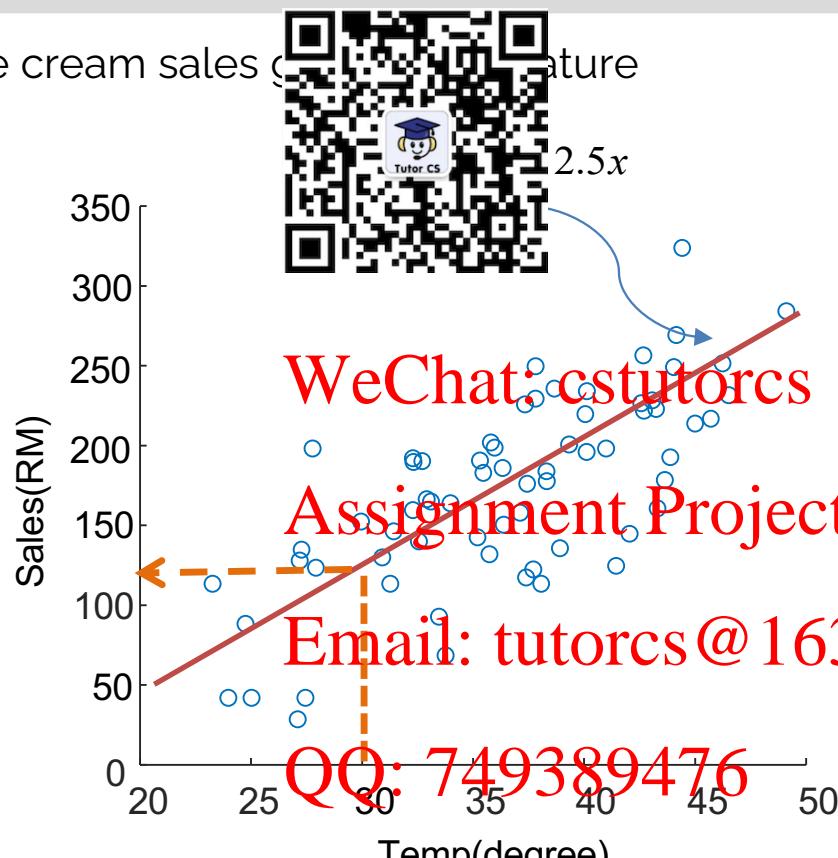
Performance metrics to assess
model accuracy

Illustration: Linear Regression model

Problem: Predict ice cream sales given temperature

Data

Day	Temp	Sales
i	x_i	y_i
1	36	200
2	31	100
3	24	50
:	:	
100	38	250



Predictive Model:

- What is good model $f(\cdot)$ to maps x to y ?

$$f(x) = \theta_0 + \theta_1 x$$

Model Learning/Estimation:

- How to choose parameters θ_0, θ_1 ?

- Define **loss function**
- Estimate using **learning algorithm**

Estimated parameters: $\theta_0 = 50, \theta_1 = 2.5$

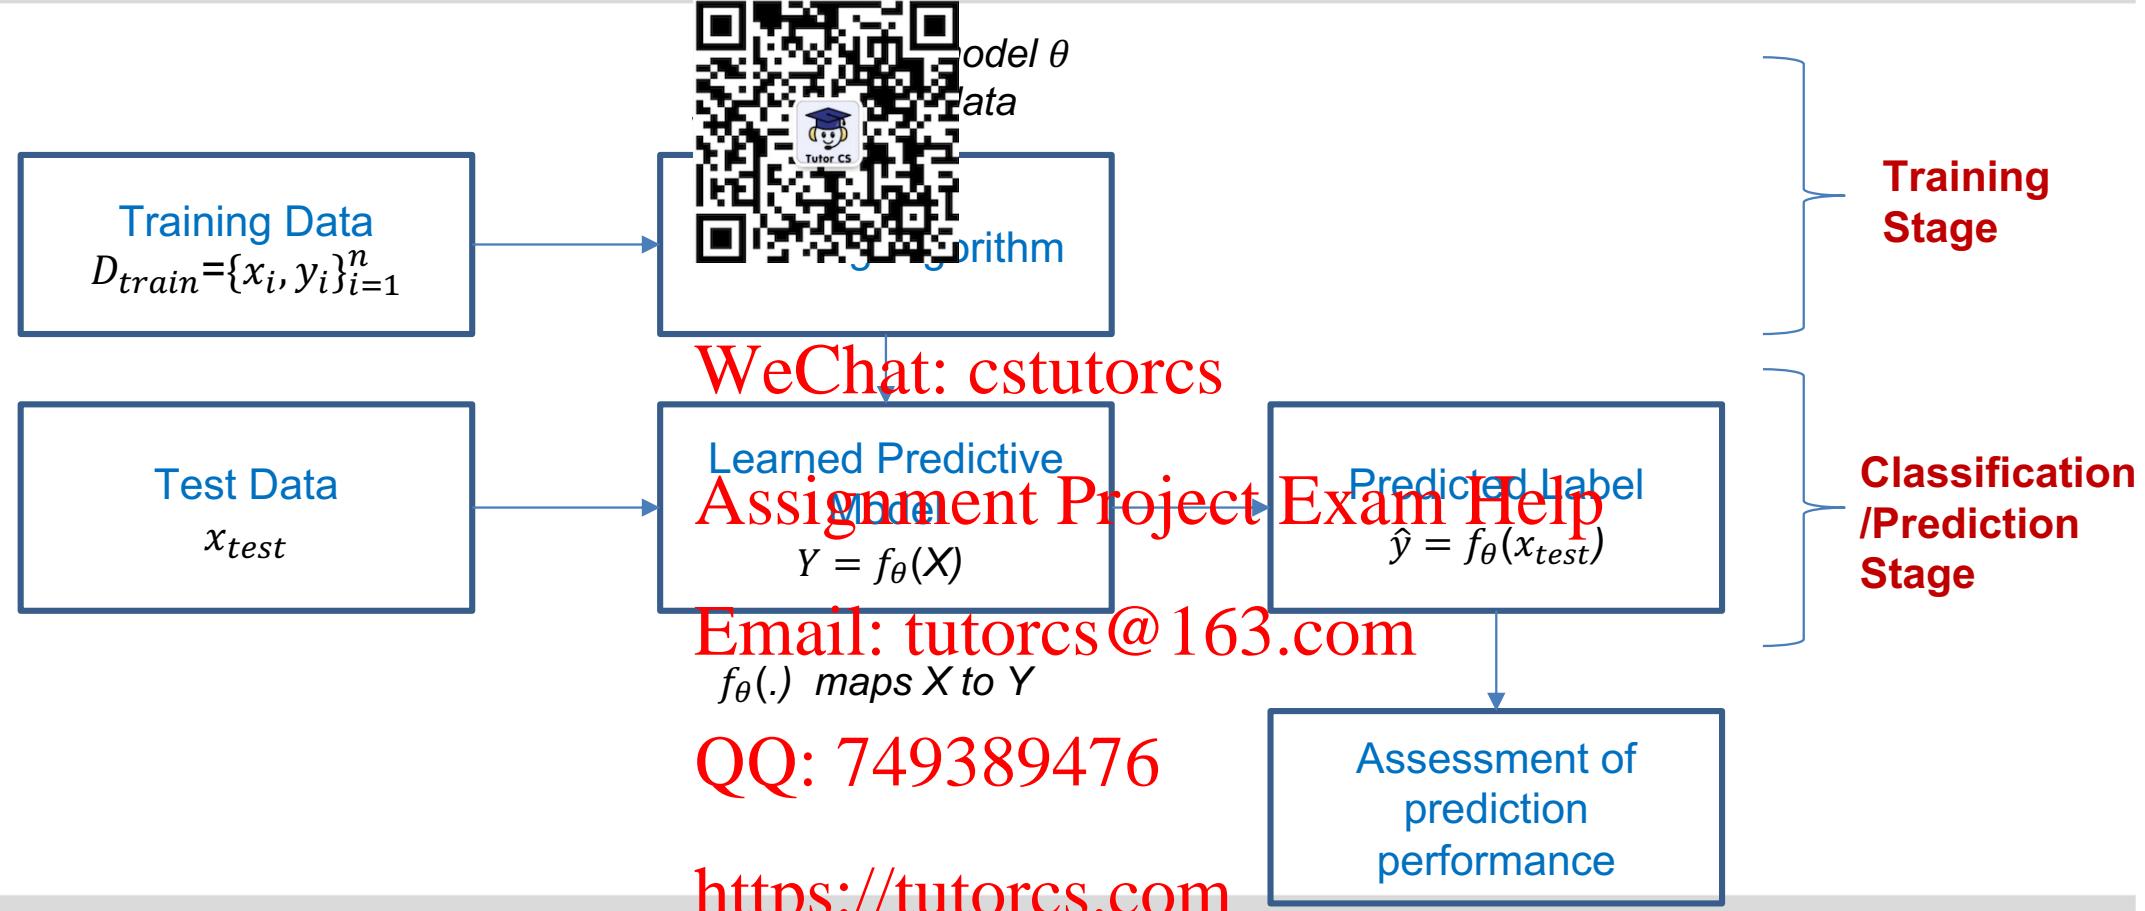
Prediction:

- Given new input, predict y with learned model $\hat{y} = f(x_{new}) = \theta_0 + \theta_1 x_{new}$

Predicted output

$$\hat{y} = 50 + 2.5(30) = 125$$

Overview of machine learning



Data

程序代写代做 CS编程辅导

Features: x_i

- a set of attributes, each is vector or matrix.
- E.g., represent each email bag-of-word vector (feature) into a real-valued matrix.



Labels: y_i

- values, categories, classes, assigned to data points.
- E.g., 0 = non-spam, 1 = spam,

WeChat: cstutorcs

Assignment Project Exam Help

Data points (aka instances, samples): $\{x_i\}$ or $\{x_i, y_i\}$

- these are items or instances of data used for training and evaluating ML models.
- E.g., labelled emails in spam detection; transaction data in credit card fraud detection; a photo in face recognition.

QQ: 749389476

<https://tutorcs.com>

data points ... $\{x_i\}$



data points with labels ... $\{x_i, y_i\}$



x_i features

y_i 0 = Jack, 1 = John, etc ...

Dataset with n samples: $\mathcal{D} = \{x_i, y_i\}_{i=1}^n$



Machine Learning: Data Types



Vector

- A mathematical
- *dense vectors*, where every entry is stored, and
- *sparse vectors*, where only the nonzero entries are stored to save space.

WeChat: cstutorcs

Labeled Point

Assignment Project Exam Help

- A labeled data point for supervised learning algorithms such as classification and regression
- Includes a feature vector and a label (which is a floating point value).

Email: tutorcs@163.com
QQ: 749389476

<https://tutorcs.com>

Machine Learning: Data Types

Vector

- A mathematical
- *dense vectors*, where every entry is stored, and
- *sparse vectors*, where only the nonzero entries are stored to save space.



Dense Vector (1.0, 0.0, 5.4, 0.0)

1.0	0.0	5.4	0.0
-----	-----	-----	-----

WeChat: cstutorcs
Assignment Project Exam Help

Sparse Vector (4, [0, 2], [1.0, 5.4])

QQ: 749389476

<https://tutorcs.com>

Email: tutorcs@163.com

1.0

5.4

1.0

5.4

1.0	0.0	5.4	0.0
-----	-----	-----	-----

Size = 4

Features

程序代写代做 CS编程辅导



All learning algorithms involve defining a set of *features* for each item, which will be fed into the learning function.

- For example, for a search, some features might include the server it comes from, or the number of mentions of the word free, or the color of the text.

WeChat: cstutorcs

In many cases, defining the right features is the most challenging part of using machine learning.

- For example, in a product recommendation task, simply adding another feature (e.g., realizing that which book you should recommend to a user might also depend on which movies she's watched) could give a large improvement in results.

<https://tutorcs.com>

Machine Learning Fundamentals



Supervised and Unsupervised Models

Bias and Variance

WeChat: cstutorcs
Underfitting and Overfitting

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Model: Types of Model Learning



Types of Model Learning: Supervised



- Goal: Learn a function from labelled training data to predict the output label(s) given unlabeled input.

WeChat: cstutorcs

- Training data consists of input features and output information (labels)

Assignment Project Exam Help

Data: $(x_1, y_1), \dots, (x_n, y_n)$

- Two types of supervised learning:

Classification

Email: tutorcs@163.com

Regression

QQ: 749389476

x = feature

y = a discrete label (classification),

y = a continuous value (regression)

<https://tutorcs.com>

程序代写代做 CS 编程辅导 Supervised Machine Learning: Classification

Classification problem: To solve classification problems, we want to map inputs into a discrete set of classes or labels.

- Binary classification
- Multinomial (Multi-class) classification



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476



Binary classification example: dog or not dog

<https://tutorcs.com>

Supervised Machine Learning: Classification



WeChat: cstutorcs
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

Multinomial classification example: Australian shepherd, golden retriever, or poodle

<https://tutorcs.com>

Supervised Machine Learning: Regression

- A regression problem is one where the output is a real value, such as “dollars” or “weight”.



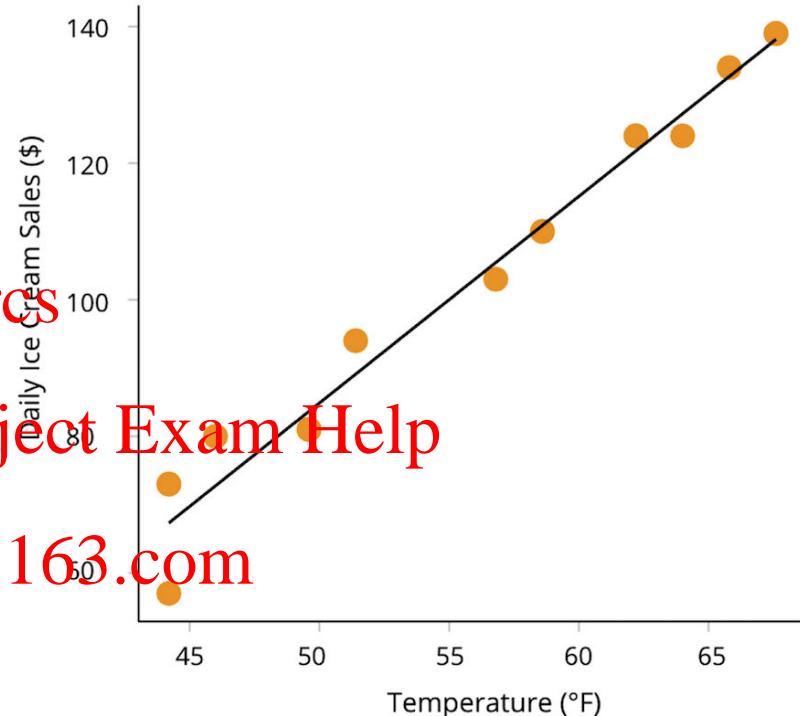
WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

Regression example: predicting ice cream sales based on temperature
<https://tutorcs.com>



Supervised Machine Learning in Apache Spark



Algorithm

Linear regression

Logistic regression

Decision trees

Gradient boosted trees

Random forests

Naive Bayes

Support vector machines (SVMs)

Practical usage

Regression

Classification (we know, it has regression in the name!)

WeChat: cstutorcs

Both

Assignment Project Exam Help

Both

Email: tutores@163.com

Classification

Classification

Types of Model Learning: **Unsupervised**

- Goal: Explore the underlying structure of the data to extract meaningful information without guidance of known output info.



- Deals with unlabelled data (no output labels)

WeChat: cstutorcs

- Two types of unsupervised learning:

- Clustering
 - Association

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Unsupervised Machine Learning: Clustering



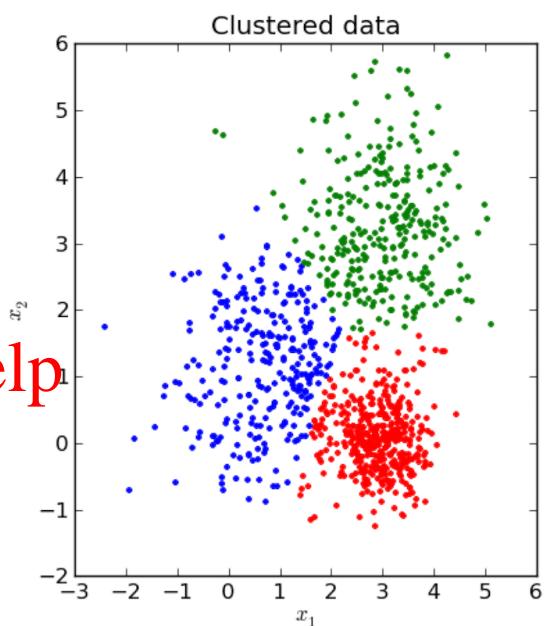
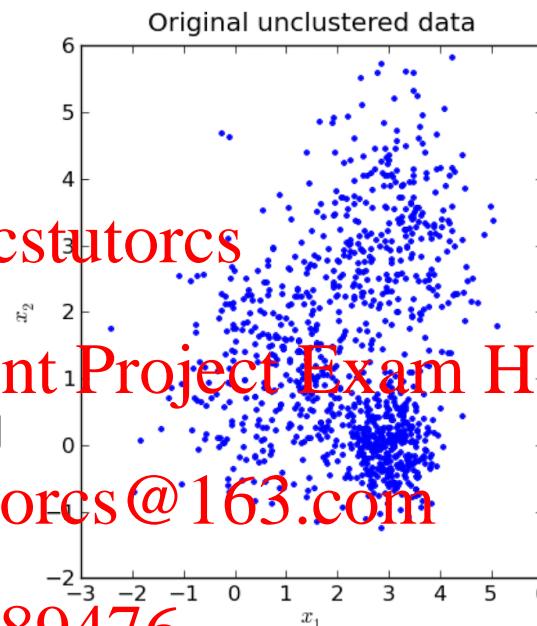
- Clustering problem: Divide data into clusters which are similar to them and are dissimilar to belonging to another cluster.

WeChat: cstutorcs

- Where you want to discover the inherent groupings in the data, e.g. grouping customers by purchasing behaviour

Email: tutorcs@163.com

QQ: 749389476



<https://tutorcs.com>

Clustering example

Unsupervised Machine Learning: Association



- Association rule learning

Discover the probability occurrence (association) between items in a large dataset

WeChat: cstutorcs

- Where you want to discover rules that describe large portions of your data, e.g., people who buy X also tend to buy Y.

QQ: 749389476

<https://tutorcs.com>

24



Unsupervised Machine Learning in Apache Spark

- k -means,
- Latent Dirichlet Allocation (LDA), and
- Gaussian mixture models.



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Machine Learning: Assessment



How to prepare the data?

- Train-Test split
- K-fold cross-validation

How to measure performance?

- TP, FP, TN, FN, confusion matrix
- Accuracy, Recall, Precision, F1-score

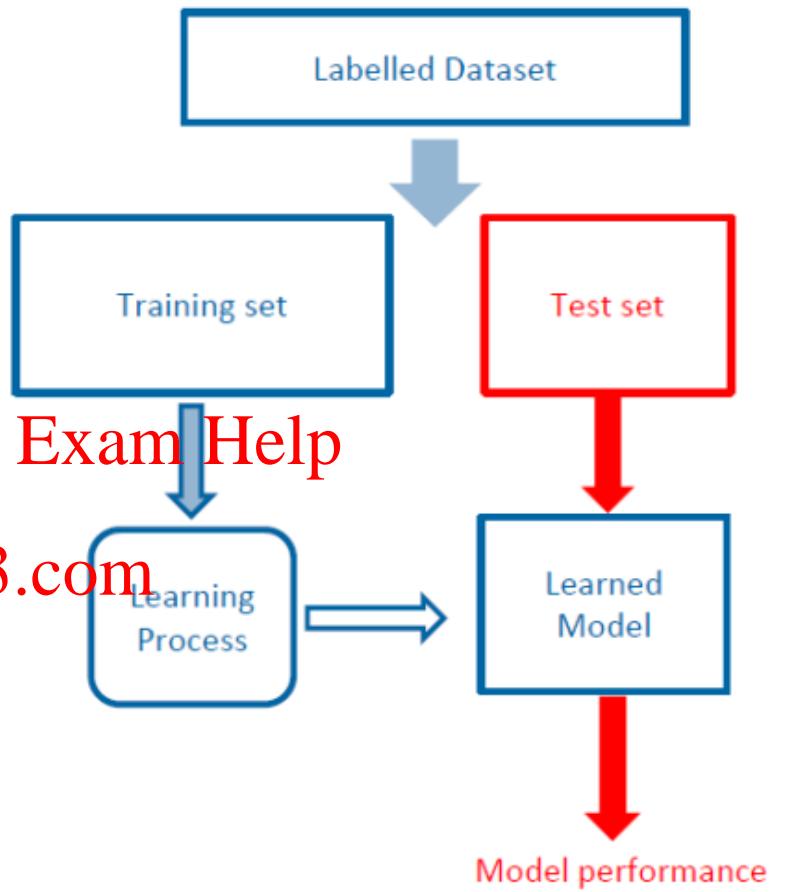
WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Machine Learning: Performance Metrics



Example: Email Spam Detection

In test set: 10 spam, 20 non-spam

Positive = spam

True labels
WeChat: cstutorcs

		SPAM (1)	NON-SPAM (0)
Predicted labels	SPAM (1)	7	5
	NON-SPAM (0)	3	15

		actual class	
		positive	negative
predicted class	positive	true positives (TP)	false positives (FP)
	negative	false negatives (FN)	true negatives (TN)

$$\text{Accuracy} = \frac{tp + tn}{tp + tn + fp + fn}$$

$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$

$$F = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Machine Learning: Bias and Variance



Bias is the gap between the averaged predicted value by the model and the actual value of the data.

Variance measures the distance of the predicted values in relation to each other.
WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

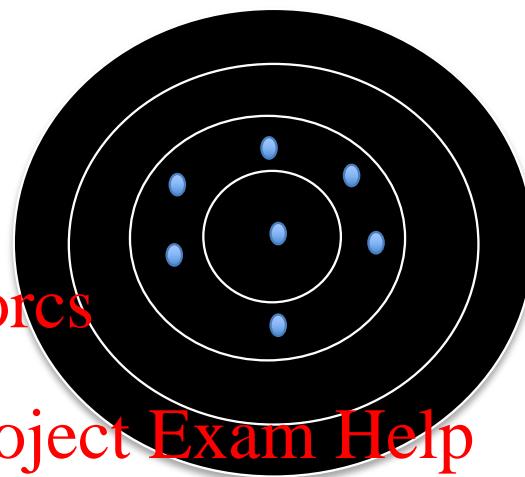
Machine Learning: Bias and Variance

程序代写代做CS编程辅导

Low Bias



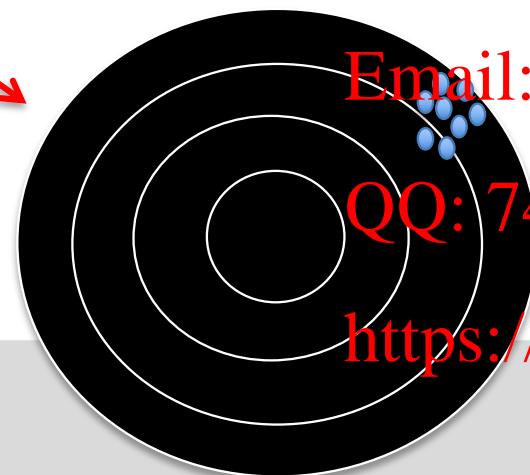
High Variance



Overfitting

Underfitting

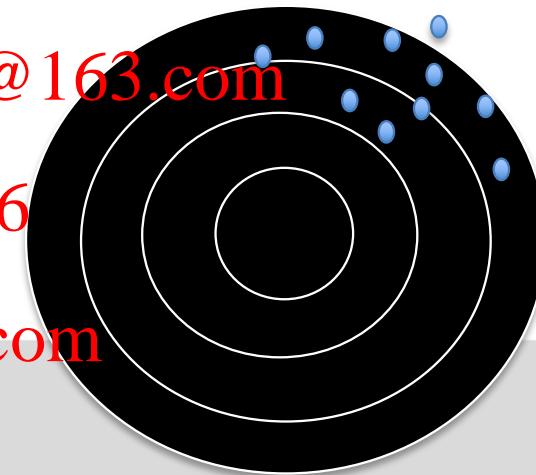
High Bias



Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Machine Learning: Overfitting and Underfitting



Overfitting (high variance, low bias) is a model that performs well on the training data but generalizes poorly to any new data.

Underfitting (low variance, high bias) is an overly simple model that does not perform well even on the training data.

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Machine Learning: Overfitting and Underfitting



- Preventing Overfitting
 - Train with more data

WeChat: cstutorcs
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



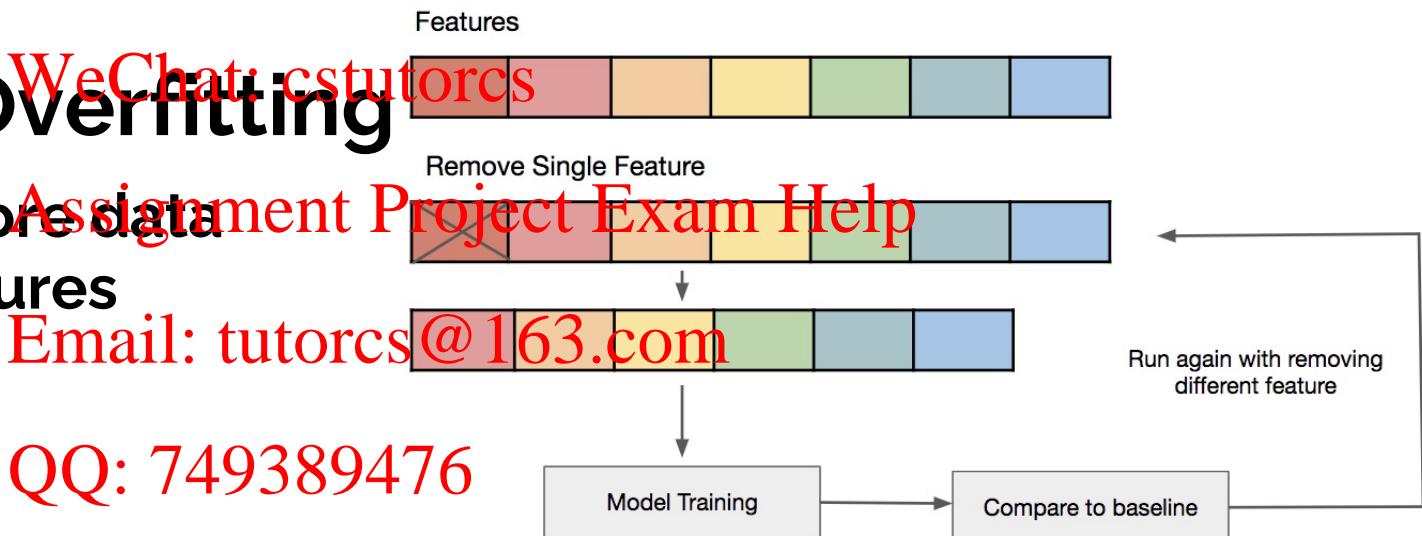
Machine Learning: Overfitting and Underfitting



Preventing Overfitting

- Train with more data
- Remove features

WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com
QQ: 749389476
<https://tutorcs.com>



Machine Learning: Overfitting and Underfitting

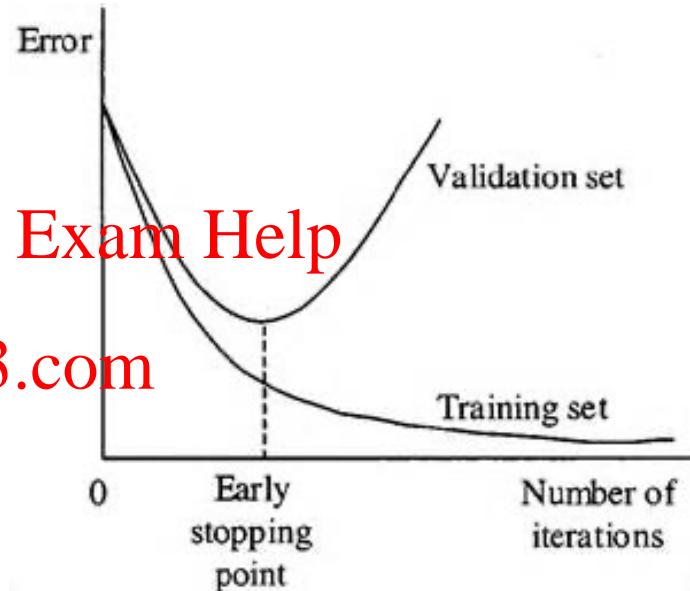


▪ Preventing Overfitting

- Train with more data
- Remove features
- Early stopping

WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com
QQ: 749389476

<https://tutorcs.com>



Machine Learning: Overfitting and Underfitting



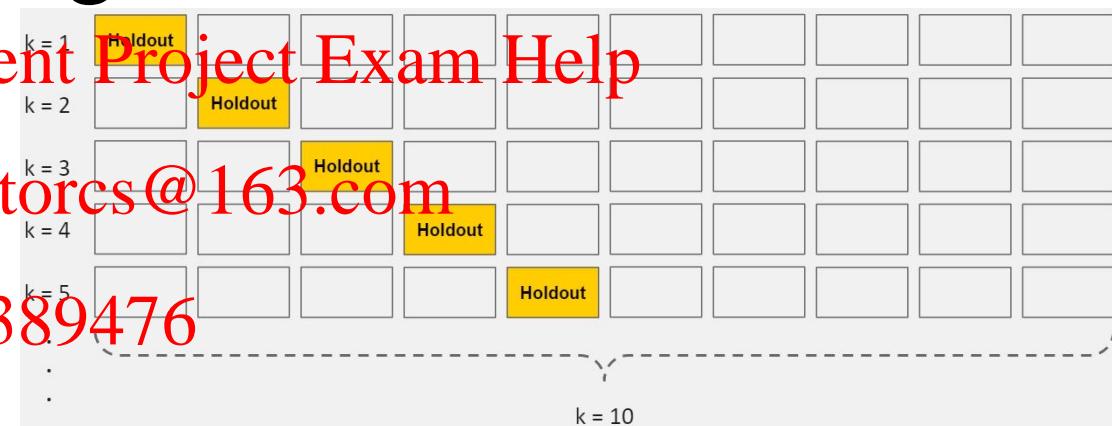
▪ Preventing Overfitting

- Train with more data
- Remove features
- Early stopping
- Cross validation

QQ: 749389476

<https://tutorcs.com>

K-Fold Cross-Validation



To be continued..

程序代写代做 CS编程辅导

Next topic -> Featurization



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



程序代写代做 CS编程辅导



MONASH
INFORMATION
TECHNOLOGY

Machine Learning- Featurization

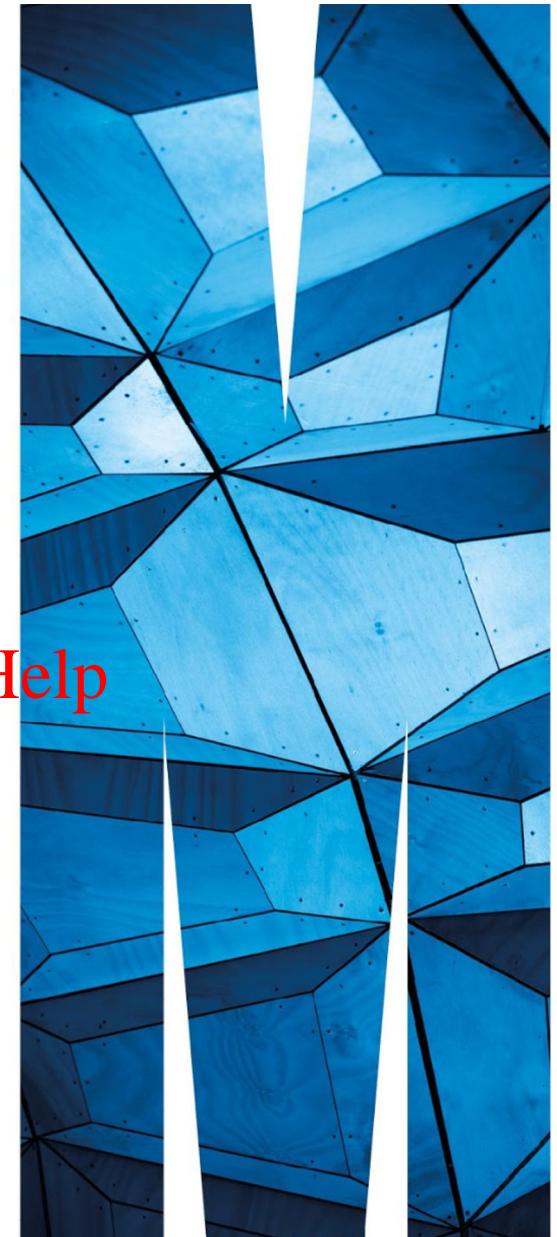
WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Machine Learning: Pipeline

程序代写代做 CS 编程辅导

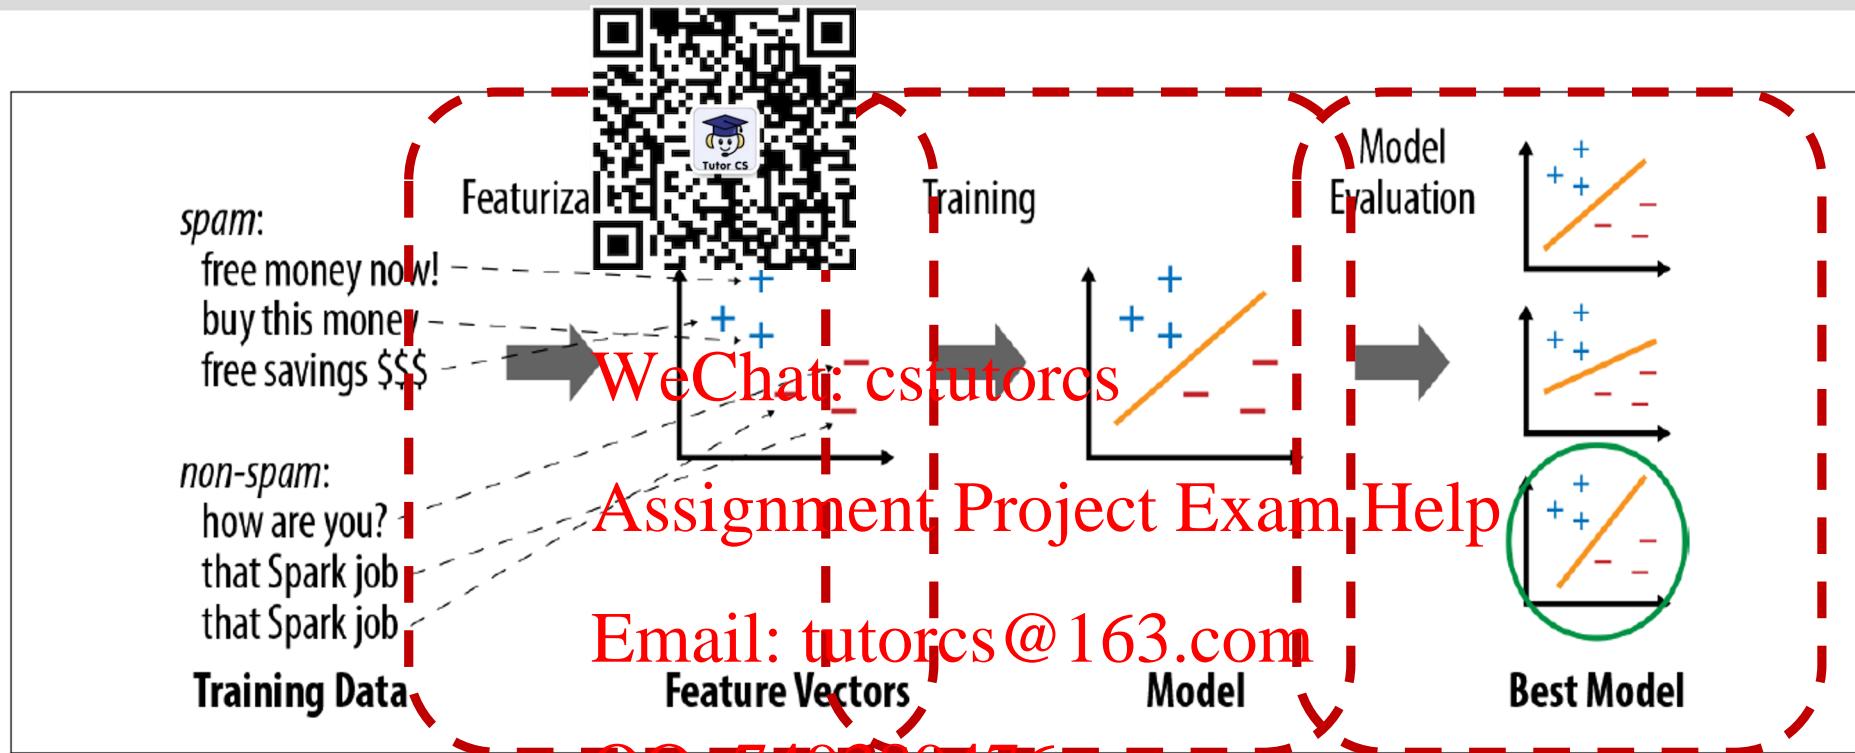


Figure 11-1. Typical steps in a machine learning pipeline

<https://tutorcs.com>

Machine Learning: Pipeline

程序代写代做 CS编程辅导

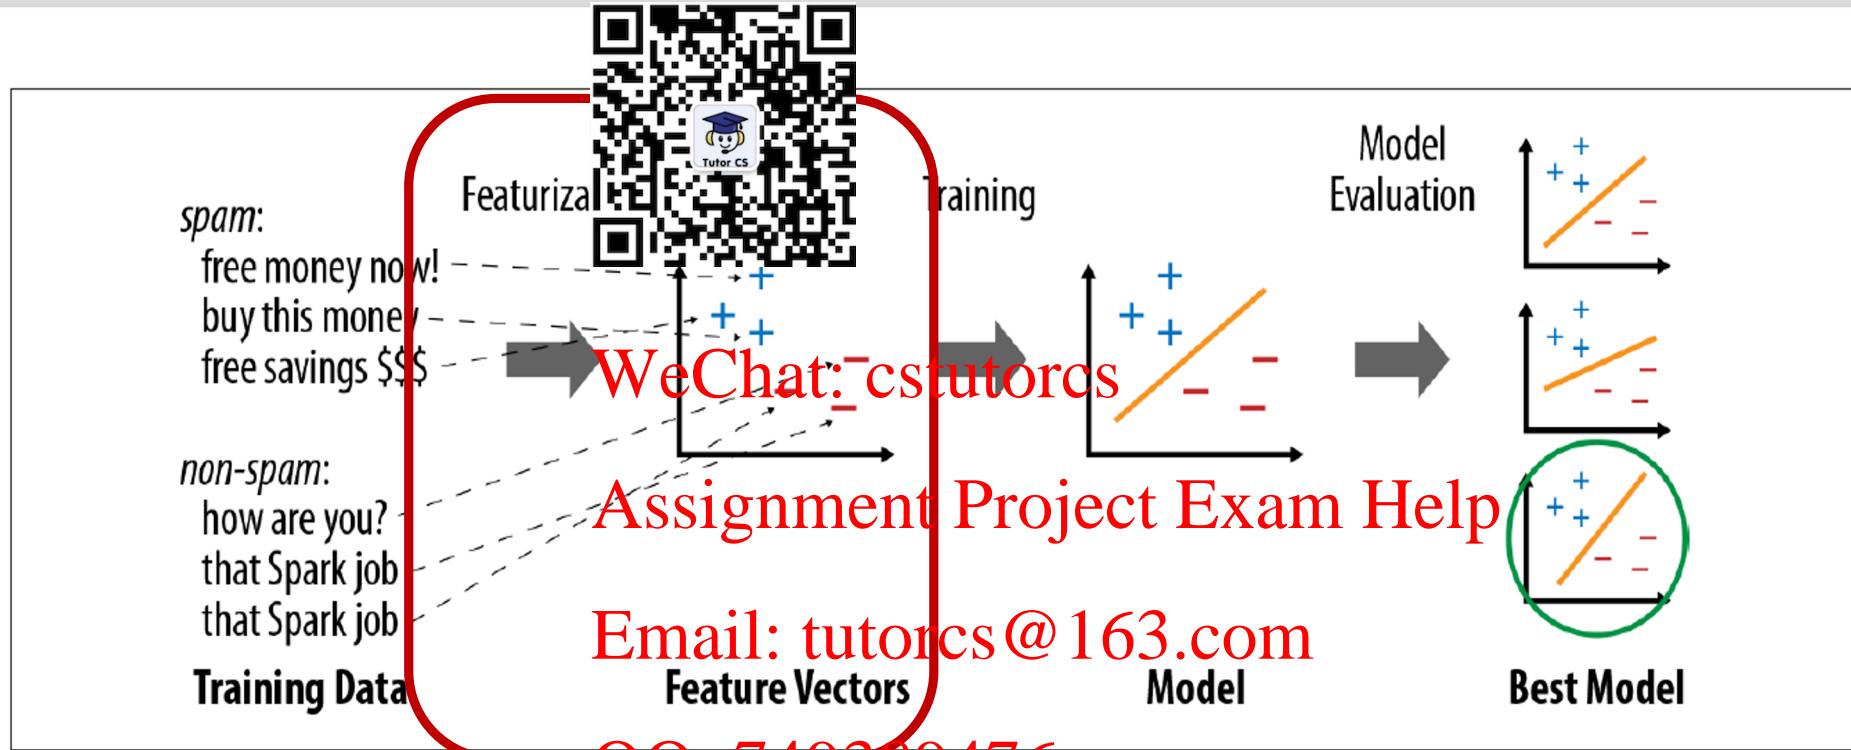


Figure 11-1. Typical steps in a machine learning pipeline

<https://tutorcs.com>

Featurization: Extraction, transformation and selection

Extraction

- Extracting features from “raw” data



Transformation

- Scaling, converting, or modifying features

Selection

- Selecting a subset from a larger set of features

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extraction and Transformation

Features

- Any machine learning algorithm requires some training data. In training data we have values for all features for all historical records. Consider this simple data set



Height	Weight	Age	Class
165	70	22	Male
160	58	22	Female

WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com

- We can prepare training data by following two techniques

Feature Extraction QQ: 749389476

Feature Selection

<https://tutorcs.com>

Featurization: Feature Extraction and Transformation



Feature extractors

- CountVectorizer
- TF-IDF
- Word2Vec
- FeatureHasher (in tutorial)

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extractors



Count Vectorizer

- Convert a collection of text documents to vectors of token counts.
- During the fitting process, Count Vectorizer will select the top `vocabSize` words ordered by term frequency across the corpus.

WeChat: cstutorcs
Assignment Project Exam Help

id	texts
---	-----
0	Array("a", "b", "c")
1	Array("a", "b", "c", "a")

<https://tutorcs.com>

Featurization: Feature Extractors



Term Frequency–I

Document Frequency, or TF-IDF,

- A simple way to extract feature vectors from text documents (e.g., web pages).
- It computes two statistics for each term in each document:

WeChat: cstutorcs

The term frequency (TF), which is the number of times the term occurs in that document and

Assignment Project Exam Help

The inverse document frequency (IDF), which measures how (in)frequently a term occurs across the whole document corpus.

Email: tutorcs@163.com
QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extractors



Term Frequency–I

Document Frequency, or TF-IDF,

- Denote a term by t , a document by d , and the corpus by D .
- Term frequency $TF(t,d)$ is the number of times that term t appears in document d , while document frequency $DF(t,D)$ is the number of documents that contains term t .

WeChat: cstutorcs
Assignment Project Exam Help

- Inverse document frequency is a numerical measure of how much information a term provides:

Email: tutorcs@163.com
QQ: 749389476

$$IDF(t,D) = \log \frac{|D| + 1}{DF(t,D) + 1},$$

where $|D|$ is the total number of documents in the corpus.

<https://tutorcs.com>

Featurization: Feature Extractors



Term Frequency–Inverse Document Frequency, or TF-IDF,

- The product of the term frequency, $TF(t, d)$, and the inverse document frequency, $IDF(t, D)$, shows how relevant a term is to a specific document ($TF \times IDF$, shows how relevant a term is to a specific document but rare in the whole corpus).
- The TF-IDF measure is simply the product of TF and IDF:

Assignment Project Exam Help

$$TFIDF(t, d, D) = TF(t, d) \cdot IDF(t, D).$$

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extractors



Term Frequency–I

Suppose that we have count tables of a corpus consisting of only two documents, as listed on the right.

Document Frequency, or TF-IDF,

Term	Term Count
this	1
is	1
a	2
sample	1

Term	Term Count
this	1
is	1
another	2
example	3

Calculate TF-IDF for the term "this".

WeChat: cstutorcs

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extractors



TF-IDF (Solution),

Calculating TF for "this"

TF ("this", d1) = 1/5 = 0.2
TF ("this", d2) = 1/7 = 0.14
(Approx.)

Document 1

Term	Term Count
this	1
is	1
a	2
sample	1

Document 2

Term	Term Count
this	1
is	1
another	2
example	3

WeChat: cstutorcs

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extractors



TF-IDF (Solution),

Calculating IDF for "th

$$|D| = 2$$

$$DF(t, D) = 2$$

$$IDF ("this", D) = \log (3/2)$$

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com
sample

$$IDF(t, D) = \log \frac{|D| + 1}{DF(t, D) + 1}$$

QQ: 749389476
<https://tutorcs.com>

13

Featurization: Feature Extractors



TF-IDF (Solution),

Calculating TF-IDF for "this"

$$\text{TF-IDF}(\text{"this"}, d_1, D) = 0.2 * \log_{10}(D / \text{Term Count})$$

$$\text{TF-IDF}(\text{"this"}, d_2, D) = 0.14 * \log_{10}(D / \text{Term Count})$$

Term	Term Count
this	1
is	1
a	2
sample	1

Term	Term Count
this	1
is	1
another	2
example	3

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extractors



Exercise: Calculate the term count for the term “example”.

Document 1

Term	Term Count
this	1
is	1
a	2
sample	1

Document 2

Term	Term Count
this	1
is	1
another	2
example	3

<https://tutorcs.com>

Featurization: Feature Extractors



Word2Vec

- maps each word to a fixed-size vector.
- transforms each document into a vector using the average of all words in the document.
- this vector can then be used as features for prediction, **Assignment Project Exam Help, document similarity calculations**, etc.

Email: tutorcs@163.com

Home Work: Do some research and write a program to find the document similarity using Word2Vec.

<https://tutorcs.com>

Featurization: Extraction, transformation and selection

Extraction

- Extracting features from “raw” data



Transformation

- Scaling, converting, or modifying features

Selection

- Selecting a subset from a larger set of features

WeChat: cstutorcs
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Extraction and Transformation



Feature Transform

- Tokenization
- Stop Words Removal
- String Indexing
- One Hot Encoding
- Vector Assembler (Implement In tutorial)

WeChat: cstutorcs
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Transformers



Tokenization

- It is the process of converting text (such as a sentence) and breaking it into individual terms (usually words).

WeChat: cstutorcs

Text
“The cat sat on the mat.”

Assignment Project Exam Help



Email: tutorcs@163.com

Tokens
“the”, “cat”, “sat”, “on”, “the”, “mat”, “.”

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Transformers



Stop Words are words that should be excluded from the input, typically because they appear frequently and don't carry as much meaning.

Some words contain more information than others
WeChat: cstutorcs

Stopwords — [the, in, for, you, will, have, be]
Assignment Project Exam Help

Quiz: How many words will be removed when we remove
stopwords from "Hi Katie the machine learning
class will be great best Sebastian"
QQ: 749389476

3

<https://tutorcs.com>

Featurization: Feature Transformers



Stop Words Removal

Takes as input a sequence of strings (e.g. the output of a Tokenizer)

Drops all the stop words from the input sequences.

WeChat: **csutorcs**

<code>id</code>	<code>raw</code>	<code>Filtered</code>
---	---	---
0	[I, saw, the, red, balloon]	[saw, red, balloon]
1	[Mary, had, a, little, lamb]	[Mary, little, lamb]

<https://tutorcs.com>

Featurization: Feature Transformers



String Indexing

Encoding a string column of labels to a column of label indices.

id	category	categoryIndex
WeChat: cstutorcs		
0	a	0.0
1	b	2.0
2	c	1.0
3	a	0.0
4	b	2.0
5	c	1.0

<https://tutorcs.com>

Featurization: Feature Transformers



One Hot Encoding

Maps a categorical feature represented as a label index to a binary vector.

A single one-value indicates the presence of a specific feature value from among the set of all feature values.

For string type input data, it is common to encode categorical features using String Indexing first.

WeChat: **csstutorcs**
Email: **tutorcs@163.com**

QQ: **749389476**

<https://tutorcs.com>

Featurization: Feature Transformers



Why One Hot Encoding?

For categorical variables when there is no ordinal relationship, the string indexing is not enough...

Using this encoding and allowing the model to assume a natural ordering between categories may result in poor performance or unexpected results.

A one-hot encoding can be applied to the integer representation. This is where the integer encoded variable is removed and a new binary variable is added for each unique integer value.

<https://tutorcs.com>

Featurization: Feature Transformers



Why One Hot Encoding?

Example: Let's say we have 3 data instances with attributes of Preferred Programming Language and OS of Choice.

Preferred Programming Language	OS of Choice
JavaScript	OSX
Python	Linux
Scala	OSX

Assignment Project Exam Help
Email: tutorcs@163.com
QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Transformers



Why One Hot Encod

String Indexing

Preferred Programming Language	OS of Choice	0	1	2	0	1	0
Javascript	OSX	0	1	0	0	1	0
Python	Linux	1	0	0	1	0	0
Scala	OSX	0	0	1	0	0	0

WeChat: csutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Transformers

- Why can't we Store Numerical Data?

The Problem Of Ordinality

Machine learning algorithms treat the ordinality of numbers in an attribute with some significance: *a higher number "must be better" than a lower number.*



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

String Indexing

Preferred Programming Language	OS of Choice
0	0
1	1
2	0

Featurization: Feature Transformers



- **String Indexing**

	Preferred Programming Language	OS of Choice
0		0
1	WeChat: cstutorcs	0
2		0

▪ **One Hot Encoding**

	Javascript	Python	Scala	OSX	Linux
0	1	0	0	1	0
1	0	1	0	0	1
2	0	0	1	1	0

Homework: Have a look at [Principal Component Analysis \(PCA\)](#)
~~QQ: 749589476~~

<https://tutorcs.com>

Featurization: Feature Transformers



Why One Hot Encoding?

For categorical variables when there is no ordinal relationship, the string indexing is not enough...

Using this encoding and allowing the model to assume a natural ordering between categories may result in poor performance or unexpected results.

A one-hot encoding can be applied to the integer representation. This is where the integer encoded variable is removed and a new binary variable is added for each unique integer value.

<https://tutorcs.com>

Featurization: Extraction, transformation and selection

Extraction

- Extracting features from “raw” data



Transformation

- Scaling, converting or modifying features

Selection

- Selecting a subset from a larger set of features

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Featurization: Feature Selectors



Feature selection

- This process tries to find the most important features that are contributing to decide the label.

Vector Slicer

WeChat: cstutorcs

- It takes a feature vector and outputs a new feature vector with a sub-array of the original features.
- It is useful for extracting features from a vector column.

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

userFeatures		features
[0.0, 10.0, 0.5]		[10.0, 0.5]

Lecture Demo

程序代写代做 CS编程辅导

Spam Classification

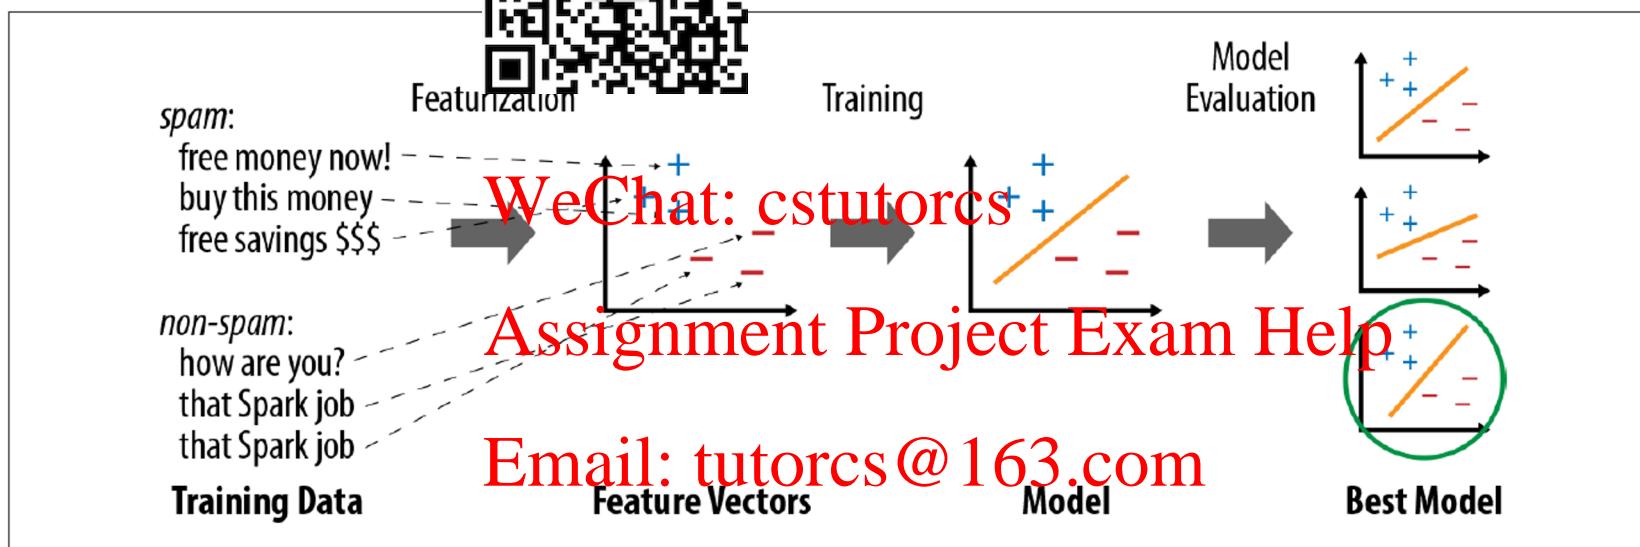


Figure 11-1. Typical steps in a machine learning pipeline
QQ: 749389476

<https://tutorcs.com>

Thank You

程序代写代做 CS编程辅导

See you next week



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>